

# Lightning Pool: A Non-Custodial Channel Lease Marketplace

Olaoluwa Osuntokun  
roasbeef@lightning.engineering

Conner Fromknecht  
conner@lightning.engineering

Wilmer Paulino  
wilmer@lightning.engineering

Oliver Gugger  
oliver@lightning.engineering

Johan Halseth  
johan@lightning.engineering

November 2nd, 2020

DRAFT (commit e30c796)

## Abstract

In this paper we examine the core inbound liquidity bootstrapping problem the Lightning Network faces, and frame it as a resource allocation problem to be solved by the application of market design and auction theory. We present Lightning Pool, a non-custodial channel lease marketplace, implemented as a sealed-bid frequent batched uniform clearing price auction which allows participants to buy and sell capital obligations on the network. We call these capital obligations *channel leases*. A channel lease can be viewed as a cross between a traditional fixed-income asset and an internet bandwidth peering agreement. Channel leases allow nodes on the network with idle capital to earn yield (based on a derived *per-block* interest rate) by selling channels to other agents in the marketplace. The duration of such contracts is enforced on-chain using Bitcoin Script. We construct **Lightning Pool** using a novel design for constructing overlay applications on top of Bitcoin: a shadow chain. Shadow chains allow users to remain in full custody of their funds at all times while also participating in higher level applications that scale via an optimistic transaction cut-through protocol.

# Contents

<b>1</b>	<b>Introduction</b>	<b>4</b>
1.1	Our Contributions . . . . .	6
<b>2</b>	<b>Background</b>	<b>6</b>
2.1	Payment Channels & the Lightning Network . . . . .	6
2.2	Liquidity Bootstrapping Problems in the Lightning Network . . . . .	8
2.3	New Routing Node Bootstrapping . . . . .	8
2.4	New Service Bootstrapping . . . . .	9
2.5	End User Bootstrapping . . . . .	10
2.6	Market Design & Auction Theory . . . . .	11
2.7	Money Markets & Capital Leases . . . . .	12
<b>3</b>	<b>Bootstrapping Problems as Solved by CLM</b>	<b>13</b>
3.1	Bootstrapping New Users via Sidecar Channels . . . . .	13
3.2	Demand Fueled Routing Node Channel Selection . . . . .	13
3.3	Bootstrapping New Services to Lightning . . . . .	14
3.4	Cross-Chain Market Maker Liquidity Sourcing . . . . .	14
3.5	Instant Lightning Wallet User On Boarding . . . . .	14
3.6	Variance Reduction in Routing Node Revenue . . . . .	15
<b>4</b>	<b>The Channel Lease Marketplace</b>	<b>15</b>
4.1	High-Level Description . . . . .	15
4.2	Lightning Channel Leases . . . . .	18
4.3	Non-Custodial Auction Accounts . . . . .	19
4.4	Order Structure & Verification . . . . .	20
4.5	Auction Design . . . . .	22
4.5.1	Auction Specification . . . . .	22
<b>5</b>	<b>The Shadowchain: A Bitcoin Overlay Application Framework</b>	<b>24</b>
5.1	High-Level Description . . . . .	25
5.2	Comparison To Related Frameworks . . . . .	26
5.3	The Shadowchain Framework . . . . .	26
5.3.1	Shadowchain Orchestrator . . . . .	26
5.3.2	Lifted UTXOs . . . . .	26
5.3.3	The Shadowchain . . . . .	27
5.3.4	Shadowchain Operation . . . . .	30
<b>6</b>	<b>Lightning Pool: A Channel Liquidity Marketplace as a Shadow Chain</b>	<b>32</b>
6.1	Instantiating a CLM . . . . .	32
6.1.1	System Initialization . . . . .	32
6.1.2	Lightning Pool Accounts . . . . .	33
6.1.3	Channel Leases in the Lightning Network . . . . .	36
6.1.4	Order Structure . . . . .	39
6.1.5	Node Rating Agencies . . . . .	40
6.1.6	Uniform Price Market Clearing & Matching . . . . .	40
6.1.7	Auction Batch Execution . . . . .	42
6.1.8	Sidecar Channel Market Clearing & Batch Execution . . . . .	45

6.1.9	LSAT as Pool Tickets . . . . .	46
6.2	The Lightning Pool Shadowchain . . . . .	46
6.2.1	Lightning Pool Accounts as Lifted UTXOs . . . . .	46
6.2.2	Auction Batch Proposal . . . . .	47
6.2.3	Shadowchain Batch Execution . . . . .	47
6.2.4	Unconfirmed Batch Cut-Through . . . . .	47
6.2.5	Auction Upgrades . . . . .	48
<b>7</b>	<b>Security Analysis</b>	<b>48</b>
<b>8</b>	<b>Future Directions</b>	<b>49</b>
<b>9</b>	<b>Related Work</b>	<b>49</b>
<b>10</b>	<b>Conclusion</b>	<b>50</b>

# 1 Introduction

**Lightning Network Bootstrapping Challenges.** The Lightning Network (LN) is the largest deployed Layer 2 payment channel network [1]. A payment channel network is comprised of a series of individual payment channels, which, when strung together, enable rapid low-latency payments between participants in the network. Due to the off-chain nature of these payments (only the final summary hits the blockchain), the cost of payments on the LN are typically much lower than an equivalent payment on the base blockchain. In order to be able to send funds on a network, a user must open a payment channel to another participant on the network. Once the channel has been opened, both participants are able to send and receive a nearly unbounded number of payments off-chain, possibly never closing the channel on-chain. Similarly, in order to receive on the network, a user requires *another individual* to open a channel *to* the receiver. A participant can only send and receive up to the total amount of Bitcoin in a channel committed to by both parties. This allocation of capital by one party to enable another party to receive funds on the network is typically referred to as *inbound liquidity* or *inbound bandwidth*. Because a would-be user of the network must somehow convince *others* to allocate capital towards them in order to receive funds on the network, the inbound bandwidth problem remains a significant barrier to the adoption and bootstrapping of the Lightning Network.

**Routing Node Capital Requirements.** In order to incentivize users to commit capital to the network so as to help other users of the network send and receive payments, each time a node forwards a payment successfully they receive a fee. As commonly implemented, this fee has a fixed base amount to be paid for all forwards independent of payment size, and a fee rate or proportional amount that must be paid for each millionth of a satoshi forwarded [22]. We refer to nodes that join in order to facilitate payments and collect forwarding fees as *routing nodes*. In order to forward a payment of size  $N$  on the network, a routing node needs to have  $N_{in} + F$  Bitcoin allocated *towards* it as inbound bandwidth, and  $N_{out}$  Bitcoin allocated to another node as outbound bandwidth. The factor of  $F$  is the fee collected by the routing node, with the following constraint being met  $F = N_{in} - N_{out}$ . Due to this requirement that there be sufficient inbound *and* outbound bandwidth, even those nodes that exist solely to help other nodes send and receive payments, themselves face the inbound bandwidth allocation problem, a fundamental bootstrapping issue for the LN.

**Market Design Resource Allocation Problems.** The field of market design is a sub-field of economics which is concerned primarily with the efficient allocation of scarce resources [16]. Within this sub-field, of interest is a branch of market design concerned with instances wherein money is used to govern the exchange of goods and services: auction design. Auction design can be used to effectively allocate scarce resources within a domain. Common established examples of market design widely used today include: carbon emissions credits, electricity markets, auctions for airport gates, and wireless spectrum auctions. In each of these examples, market design is used to allow more effective communication of pricing information, resource availability, and the expression of preferences. Our first insight is framing the solution to the inbound bandwidth bootstrapping problem within the lens of market design. In the context of the LN, the scarce resource we aim to more efficiently allocate is inbound channel

bandwidth.

**LN Bootstrapping as Resource Allocation Problem.** In the absence of a proper venue, those that need inbound capital to operate their Lightning services often turn to soliciting capital on various chat groups, forums, or public venues like Twitter. On the other side, those seeking to deploy capital in order to facilitate network operation and gain routing fees must guess as to exactly where their capital is most demanded. As node operators may not necessarily know where their capital is most demanded, they risk opening channels to locations where they aren't actually needed, leading to poor resource utilization and capital inefficiency. It's as if node operators are speculatively building roads that no one will use (why isn't my node forwarding?), and those seeking to receive aren't able to flag their service as an attractive destination to be connected to internal network "highways". Lightning Pool solves this resource allocation problem using an auction that matches up those seeking to deploy capital (by opening channels) to those that need these channels to operate their Lightning Service or business. With each executed batch, the participants of the auction derive a *per block* interest rate which is effectively the current lease rate for capital on the Lightning Network. The auctioneer or an independent agency is also able to provide Node Ratings to participants of the network, which can be used to make more informed decisions with respect to the *quality* of the channel lease being purchased.

**Channel Lease Marketplaces.** In this paper, we present Lightning Pool, a non-custodial channel lease marketplace (CLM) that draws on modern auction theory to construct an auction that enables participants to buy and sell inbound channel bandwidth. Participants in the marketplace buy and sell a channel capital obligation, which we call a Lightning Channel Lease (LCL). An LCL is similar to a traditional bond in that one party acquires capital from another for productive use, with the party parting with their capital being compensated for their cost of capital. However as the funds within an LCL can only be used in the Lightning Network for sending/receiving, an LCL is analogous to the creation of a new virtual "road" within the LN connecting two destinations. Critically, when one purchases an LCL, the period of time those funds must be committed is enforced on-chain using Bitcoin Script. As a result, buyers of inbound channel bandwidth can be sure the capital will be committed for a set period of time. The auction itself contains several sub-auctions for the exchange of particular duration intervals expressed in blocks (similar to the various U.S Treasury auctions [3]). A non-trusted auctioneer facilitates the marketplace by accepting sealed-bid orders, clearing the market using a uniform clearing price for each duration bucket, and finally executing batches of contracts using execution transactions that update all involved accounts, delivering the purchased channels to all parties in an atomic manner. Lightning Pool and the LCL solve the inbound bandwidth problem by allowing participants on the network to effectively exchange pricing signals to determine exactly *where* in the network capital should be allocated.

**Shadowchains as an Application Framework.** Lightning Pool is the first application built on top of Bitcoin that utilizes the *Shadowchain* paradigm to construct an application-specific overlay system on top of existing Bitcoin unspent transaction outputs (UTXOs). **A user joins a shadowchain by creating a special multi-signature based output using a public key provided by the target shadowchain manager.** Once a user has joined the shadowchain, proposed

state transitions are packaged up (in the form of a block) by the shadowchain manager and proposed to each active participant. A shadowchain block updates the application state of all participants and is embedded within a normal Bitcoin transaction. Depending on the application, a shadowchain block may be indistinguishable from a normal Bitcoin transaction. Shadowchains are able to *compress* state transitions *off-chain* by employing a form of multi-party transaction cut-through [14]. As participants of a shadowchain remain in custody of their funds at all times, complex fraud proofs or exit games are unnecessary, significantly simplifying the implementation of a given shadowchain. We note that the shadowchain concept itself can be used for other applications as well.

## 1.1 Our Contributions

In summary, we make the following contributions:

- We propose a solution to the inbound bandwidth problem of the LN in the form of a marketplace to buy and sell inbound bandwidth obligations.
- We propose the Lightning Channel Lease, an inbound bandwidth obligation contract that pays a per-block interest rate to the seller from the buyer, and whose duration is enforced on-chain with Bitcoin Script.
- We put forth the concept of a Node Rating agency for channel leases in order to provide marketplace participants with information about the *quality* of a channel lease.
- We construct a new system, Lightning Pool, which is a non-custodial marketplace with off-chain order submission and on-chain batch execution that allows parties to exchange LCLs in an atomic manner.
- We design a new Bitcoin application design framework, the shadowchain, which can also serve other applications.

## 2 Background

In this section, we aim to introduce some necessary background that will be built upon in later sections to construct our solution. First, we’ll describe multi-hop payment channels and the Lightning Network as deployed today. Next, we’ll explore the nature of the inbound bandwidth bootstrapping problem the Lightning Network faces today. Along the way, we’ll explain the dynamics of routing nodes in the network, as they’re a key component of the system. Next, we introduce the field of market design and more specifically the sub-field of auction design, to demonstrate how auction design can be used to solve resource allocation problems in the real world. Finally, we provide some brief background on money markets in the traditional financial system, and how this relates to our concept of channel leases.

### 2.1 Payment Channels & the Lightning Network

#### Basic Payment Channels

A payment channel in its simplest form is an on-chain 2-of-2 multi-signature output created by parties  $A$  and  $B$ . One or both parties deposit funds into a Bitcoin Script output constructed using two public keys  $P_a$  and  $P_b$ . The transaction that creates this multi-sig output is referred to as the *funding transaction*. Before broadcasting the funding transaction, another transaction dubbed the *commitment transaction* is constructed using a series of agreed upon parameters by the two parties [21]. The commitment transaction spends the funding transaction and creates two new outputs  $D_a$  and  $D_b$  that if broadcast, will *deliver* the up-to-date balances allocated between the parties to the channel. Once the funding transaction is confirmed and broadcasted, both parties are able to rapidly update the balance of the delivery outputs,  $D_a$  and  $D_b$  so as to facilitate efficient payments between the parties.

### Bi-Directional Payment Channels

In order to safely make bi-directional payments between both parties to a payment channel, modern channel designs also employ a *commitment invalidation mechanism* [2] to ensure that only the latest commitment transaction state can be broadcasted and redeemed via the underlying blockchain. The most commonly used commitment invalidation scheme is the **replace-by-revocation** construct. In this construction, during channel negotiation, a security parameter  $T$  (which may be asymmetric for both parties) is negotiated. Using this value  $T$  which is typically expressed in blocks, a commitment transaction state can only be fully redeemed by the broadcasting party after a period of  $T$  blocks has passed. During this interval, the non-broadcasting party  $P_{defender}$  is able to provide the contested delivery output  $D_{a_i}$  with a valid witness  $W_{r_n}$  which *proves* that there exists a *newer* state  $n$  with  $n > i$  that has been ratified by both parties. The exact details of this construct are outside the scope of this paper, but Bitcoin Script and basic cryptography are used to allow a defending party to present an objective statement of contract violation by the opposing party.

### Hash Time Lock Contracts & Multi-Hop Payments

The final component of modern multi-hop payment channels is the Hash Time Locked Contract, or the HTLC. The HTLC enables payments to travel over a *series* of payment channels, allowing a set of interlinked payment channels to be composed into a logical *payment network*. An HTLC can be viewed as a specific case of a time locked commit and reveal puzzle. Loosely, an HTLC consists of four parameters: the public key of the sender  $P_s$ , the public key of the receiver  $P_r$ , the payment amount expressed in satoshis  $A_{sat}$ , a payment secret  $r$  s.t  $H(r) = h$ , and an absolute block timeout  $T$ . Given these parameters, a Bitcoin Script is set up such that, the funds deposited in the script hash output can be redeemed by the receiver  $P_r$  via a public key signature by their public key and the revelation of the payment pre-image  $r$ , or by the sender  $P_s$  after the absolute timeout  $T$  has elapsed. This construct can be chained by several parties (up to 20 in the modern Lightning Network [20]) to create a multi-hop payment within the network. One implication of this security model is that each party must ensure that their outgoing hash lock puzzle’s absolute timelock  $T_o$

is offset from the incoming absolute timelock  $T_i$  by a value of  $C_{delta}$ . This value  $C_{delta}$  is commonly referred to as the CLTV delta [22]. This value  $C_{delta}$  is an important security parameter, as if  $C_{delta}$  blocks passes and the outgoing hash lock isn't fully resolved, then a *race condition* occurs as the time out clauses of *both* the incoming and outgoing hash locks have expired.

### Routing Nodes as Profit-Seeking Capital Allocators

Entities on the Lightning Network that exist primarily to collect fees for successfully forwarding payments are referred to as *routing nodes*. A routing node commits capital to the network within payment channels in order to be able to facilitate payments in the network. As routing nodes incur an opportunity cost by committing capital to the network, they specify a fee  $F$  to be paid upon completion of a successful payment forward. This fee  $F = F_{base} + F_{rate} \cdot A_{sat}$  is comprised of two parts: a proportional amount (a rate) and a fixed amount, which are both expressed in *millisatoshis*, which are 1/1000 of the base satoshi unit.

Note that routing nodes are not compensated on an ongoing basis, and are not compensated for anything other than a completed payment. As a result, many routing nodes may be allocating capital in a non-productive manner [5] as they've speculatively opened channels to areas of the network where no true transaction demand exists. If the Lightning Network was a physical transportation network, then it would be as if eager contractors started building roads to seemingly random destinations, only to find that those roads weren't actually demanded at all. This information asymmetry (where new channels are actually demanded) and the current inability for today's network participants to exchange these key demand signals lies at the crux of the bootstrapping problems of the Lightning Network.

## 2.2 Liquidity Bootstrapping Problems in the Lightning Network

In this section, building on the background provided above, we aim to detail the various liquidity bootstrapping problems that exist in the Lightning Network today. These problems will serve motivation for our solution, the Channel Lease Marketplace, and a specific instantiation of such a construct: Lightning Pool.

## 2.3 New Routing Node Bootstrapping

As the Lightning Network is a fully collateralized network, in order to *join* the system, a participant must commit capital in the form of Bitcoin charged into payment channels on the network. Routing nodes, however, are in a unique situation, as they need to both *commit* their own capital to the network, as well as *solicit* committed capital from *other* routing nodes. This is due to the fact that in order to be able to forward a payment of size  $P_{sat}$ , the routing node must first have  $P_{sat_{out}}$  satoshis committed as *outbound* payment bandwidth (to use for sending) and  $P_{sat_{in}}$  committed as *inbound* payment bandwidth, with the difference of the two amounts,  $F = P_{sat_{in}} - P_{sat_{out}}$  being collected as a forwarding fee upon payment completion. This *pair-wise* capital commitment requirement



is commonly cited as a major barrier to Lightning Network adoption, as well as why large "hubs" are inherently economically inefficient.

A routing node operator faces two key questions when attempting to join the network in a productive manner, while also attempting to optimize for *capital* efficiency:

1. *Where* should I open channels (thereby committing outbound capital) within the network in order to *maximize* the velocity of transactions through my channels, along with the corresponding fee revenue  $F_r$ ?
2. *How* can I attract *other* routing node operators to commit capital to my node such that I can actually forward payments to earn *any* revenue  $F_r$ ?

We argue that the above two questions, optimizing for capital efficiency and velocity of committed channels, can only properly be addressed by the *existence* of a *marketplace* that allows agents (routing node operators) to communicate their preferences using demand signals. Intuitively, a channel open to an undesirable location (possibly over-served) will have low transaction velocity  $C_v$ , and result in overall lower total fee revenue  $F_r$ . In order to maximize both  $C_v$  and  $F_r$ , a routing node should only open channels to where they're *most demanded*. If an agent is willing to pay up to  $P_{premium}$  Bitcoin for inbound bandwidth, then they must gain more utility than the paid premium  $P_{premium}$ , as otherwise, such a transaction would not be economically rational. Thus, the existence of a marketplace that allows routing nodes to efficiently commit their outbound capital, as well as *purchase* new inbound capital is a key component to solving the bootstrapping problem for routing nodes.

## 2.4 New Service Bootstrapping

If routing nodes are the backbone or highway of the Lightning Network, then so called Lightning Services are the primary *destinations* for a given payment. For simplicity, we assume that a given Lightning Service is primarily a payment *sink*, in that it's primarily *receiving* over the LN. Eventually, it may become common for a service to be balanced in terms of sending *and* receiving, resulting in a net-flow of zero, but today in the network, most flows are uni-directional, creating the need for on/off chain bridges such as Lightning Loop.

### Demand for Incoming Bandwidth

Focusing on the case of a Lightning Service that's primarily a *payment sink*, in order to receive up to  $N$  Bitcoin, the service requires  $S_b$  Bitcoin to be committed as inbound capital, with  $S_b > N$ . Otherwise, assuming only channel churn, all inbound bandwidth will become saturated, rendering a service unable to receive additional Bitcoin over the LN. Therefore, the operative question a service operator needs to ask when attempting to join the network is:

- How can I solicit enough inbound bandwidth to be able to receive up to  $S_b$  Bitcoin?

### Preference for Quality of Bandwidth

It's important to note, that as operating a *valid* routing node on the network requires a degree of skill and commitment, some routing nodes are able to provide more *effective* service than others. As an example, imagine a routing node **Bob**, who has sufficient capital committed to his node in both the inbound and outbound directions, but who is chronically *offline*. As a node must be *online* in order to be able to forward payments, any capital committed by **Bob**, can essentially be considered dead weight. With this insight in mind, we revisit the bootstrapping questions of the Lightning Service to also require a high *quality of service*:

- How can I solicit enough *high quality* inbound bandwidth within the network to be able to receive up to  $S_b$  Bitcoin?

### Time Committed Incoming Bandwidth

However, from the point of view of an active Lightning Service, just having sufficient high quality inbound bandwidth may not be enough. Consider that a high quality node *Carol* may erroneously decide to commit capital elsewhere, resulting in overall lower channel velocity  $C_v$  for their channels. This type of fair-weather behavior serves as a detriment to our Lightning Services; they're unable to properly plan for the future, as they don't know *how long* the inbound bandwidth will be available for receiving payments. As a result, it's critical that the Lightning Service has a hard guarantee with respect to *how long* capital will be committed to their node. Taking this new criteria into account, we further revisit our new service bootstrapping problem statement:

- How can I solicit enough *high quality* inbound bandwidth to be able to receive up to  $S_b$  Bitcoin, that will be committed for *at least* time  $T_{blocks}$ ?

Summarizing, in addition to the existence of a marketplace for buying and selling capital commitment obligations, a would-be buyer requires some sort of *rating system* to reduce information asymmetry (distinguish the good nodes from the lemons), and also requires that any capital committed must be committed for a period of  $T_{blocks}$ . These new requirements argue for the existence of a Node Rating agency, as well as a facility that ensures that capital will be committed for a set period of time in a trust-minimized manner.

## 2.5 End User Bootstrapping

Finally, we turn to the end users of the system. In our model, the end users of the system are those that are frequently *transacting*. If routing nodes are the highways in our payment transportation network, with Lightning services as popular destinations, then users trigger *payment flows* that traverse the backbone created by routing nodes, to arrive at the Lightning services. Note that within our model, we permit end users to both send *and* receive. Compared to bootstrapping a new user to a Layer 1 system such as the Bitcoin blockchain, bootstrapping to a Layer 2 system like the Lightning Network presents additional challenges. The core challenge is created by the constraint that in order for a user to send  $K_s$  Bitcoin, they also need  $K_s$  Bitcoin committed within the network. Similarly, in order to receive up to  $K_r$  Bitcoin, they need up to  $K_r$  Bitcoin committed as inbound bandwidth.

From the perspective of attempting to achieve a similar user-experience as a base Layer 1 system, the receiving constraint is the most challenging. Notice that a user cannot simply download a Lightning wallet and start receiving funds. Instead, they need to first solicit *inbound* capital to their node first. Many wallet providers such as Breez and Phoenix have started to overcome this issue by committing capital to the users themselves. This is essentially a customer acquisition cost: by providing this inbound bandwidth to users, the wallet becomes more attractive as it enables both sending and receiving. However, just receiving isn't enough. A user needs to be able to send *and* receive. In addition to this required symmetry, a typical user also has all the same quality of service, and time-committed capital requirements as well.

With this background, we can phrase the end user bootstrapping problem as follows:

- How can a new user join the Lightning Network in a manner that allows them to both send and receive to relevant destinations in the network?

## 2.6 Market Design & Auction Theory

In this section, we make a brief detour to the economic field of auction design to examine how similar resource allocation problems have been addressed by market design in existing industries. These examples include both digital and physical goods. In the modern age, market design and proper construction of corresponding *auctions* can be used to improve resource utilization and capital efficiency [16]. Within a particular domain, context-specific design decisions can be made so as to better optimize resource allocations for all participants. Common uses of auction design in the wild include wireless spectrum auctions by the Federal Communications Commission (FCC), package auctions for auctioning off takeoff and landing rights at airports, real-time electricity markets, and also carbon credits. Market design bridges both theory and practice in order to solve real-world resource allocation constraints [15].

A commonly used tool in the field of market design is the concept of *auctions*. Auctions allow agents to gather and exchange pricing signals in order to determine who gets which goods, and at which price. The design of a proper auction for a particular resource allocation problem has a vast design space. For example, should a first or second price auction be used [11]? How frequently should the auction run? What *type* of auction should be used? Should participants be able to see the bids of other participants? And so on.

Building off the series of bootstrapping problems posed above, we turn to market design as a tool to efficiently allocate our scarce resource in question: inbound channel bandwidth. Our problem-space is unique, however, in that as we're dealing with the allocation of capital, there are inherent opportunity costs: why should a routing node commit capital to the Lightning Network, compared to some other asset that has a similar risk adjusted rate of return? In this context, our end solution may take the form of a *money market*, which is used by entities to trade short-term debt instruments.

## 2.7 Money Markets & Capital Leases

In traditional financial markets, money markets are used to allow entities to trade short term debt instruments. Examples of such instruments include U.S Treasury Bills, certificates of deposit, and repurchase agreements. Capital markets on the other hand, are the long-term analogue of money markets, in that they deal with longer timeframes, and also are more heavily traded on secondary markets with retail traders being more involved.

In the context of the Lightning Network, our concept of capital obligations appears similar to a bond, in that we require a period of time for which capital is allocated. However, unlike a bond, the committed funds can only be used on the Lightning Network to provide a new type of service: the *propensity* to receive or send funds on the network. As a result, we don't require funds in channels to be borrowed, instead they only need to be *leased* for a period of time. Also unlike bonds, wherein it's possible for the issuer of a bond to *default*, thereby failing to repay the borrowed money, in the context of the LN, there is *no inherent default risk*. Instead, arguably the concept of channel leases can be viewed as a risk free rate of return in the context of Bitcoin, and specifically in the context of the Lightning Network.

The existence of a *channel lease* serves to provide routing nodes with an additional monetary incentive (in the form of a premium paid by the lessee of the coins) to operate a routing node. As a result, we can model the revenue  $R_c$  earned by a routing node for a given channel  $C$ , as a function of the lease interval  $T$  and channel size  $A_{sat}$ . Factoring in transaction values during the interval (as routing fee revenue is a function of them), we assume that transaction values fall in the range of:  $[1, A_{sat}]$  satoshis and follow a distribution  $K$ . Given these considerations, we express the revenue of a given routing node as:

$$R_c(T, A_{sat}) = (P_c \cdot T \cdot A_{sat}) + \sum_{k=1}^{A_{sat}} \Pr[K = k] \int_{i=0}^T F_c(k) \cdot X_t(k, C) dt \quad (1)$$

where  $P_c$  is the current *per-block* interest rate,  $(F_c(k) \cdot X_t(C))$  is the *expected* routing fee revenue of the channel within that interval, for a payment of  $k$  satoshis, and  $X_t(k, C)$  is a function describing the random event of a payment of  $k$  satoshis passing through the channel  $C$  on the outgoing edge during a time-slice  $t$ .

Our model is similar to the one put forth in [32], however were concerned with the fee revenue over an interval rather than the gain of an objective function. We reference an *expectation* for fee revenue, as fees are effectively a *speculative* component of the routing revenue of a node. If a channel was allocated to a node in high demand, one would expect the latter portion of the question to possibly dominate the premium. If the opposite is the case, then a routing node would derive most of its revenue from the yield earned by leasing a channel. In this manner, the existence of a concept such as a channel lease actually serves to *reduce* the variance in a routing node's revenue, similar to how joining a mining pool can reduce the variance of a Bitcoin miner's earnings [12].

Finally, we argue that the existence of a channel lease that pays a premium based on a *per-block* interest rate would result in a novel low-risk yield-generating instrument for the greater Bitcoin network. Such a per-block interest

rate  $r_{b_i}$  would serve to allow market participants to effectively price the cost of capital on the Lightning Network. Assuming the existence of varying durations  $D_1, \dots, D_n$ , a *yield* curve conveying the relative short and long term interest rates of channel yields could be constructed. Such an instrument would then potentially serve as the basis for higher level structured products and derivatives built on top of the base channel lease instrument.

### 3 Bootstrapping Problems as Solved by CLM

Prior to outlining our design for a channel lease marketplace, we seek to provide a set of real-world cases that demonstrate the benefit of such markets for the Lightning Network.

#### 3.1 Bootstrapping New Users via Sidecar Channels

A common question concerning the Lightning Network goes something like: Alice is new to Bitcoin entirely, how can she join the Lightning Network without making any new *on-chain* Bitcoin transactions? On-boarding for a non-Lightning Bitcoin user is as simple as sending coins to a fresh address. For off-chain payment channel networks to achieve widespread usage, a similar, seamless on-boarding flow should exist.

We frame the solution to this use case in our model of channel liquidity markets. In this case, Alice is a *new* user to the network that requires *inbound* and *outbound* liquidity. Without outbound liquidity, she's unable to send to any other node on the network. Without inbound liquidity, she's unable to receive payments via the network. "Sidecar channels" allow an acquaintance of Alice, let's call her Carol, to engage in a protocol with an existing routing node on the network, Bob, to provide *both* inbound and outbound liquidity for Alice. Carol is able to provide liquidity with an off-chain, or on-chain payment. At the end of the engagement, Carol has provided channel liquidity to Alice via Bob, who himself is compensated accordingly for his role in the protocol.

#### 3.2 Demand Fueled Routing Node Channel Selection

A "routing" node on the Lightning Network is a node categorized as having a persistent, publicly reachable Internet address, a set of inbound channels from leaf nodes, and an intention to actively facilitate payment flows on the network in return for fee revenue. A common question asked in the initial bootstrapping phase of the Lightning Network by node operators is: "where should I open my channels to, such that they'll actually be routed through"? We posit that channel liquidity markets provide the answer to this question.

Channel liquidity markets can be combined with autopilot [?] techniques that automatically manage channel creation based on static and dynamic graph signals. A key drawback of autopilot techniques alone is that for the most part, they're devoid of *economic* context. A particular location in the sub-graph may be "fit" or attractive from a graph theoretic perspective, but may not lead to a high velocity channel, as there might not be inherent demand for a channel created at that particular location. Using a CLM, a node operator can enter a targeted venue to determine what the *time value* of his liquidity is on the

network. New services such as exchanges or merchants on the network can bid for the node's liquidity in order to serve their prospective customers, with the node earning a small interest rate up-front for committing his liquidity in the first place (scaled by the worst-case CSV delay).

### 3.3 Bootstrapping New Services to Lightning

Any new service operator or merchant that wishes join the Lightning Network faces the same problem: "How can I incentivize nodes to create *inbound* channels to my node in order to be able to accept payments?". CLMs provide an elegant solution. The merchant/exchange/service uses their existing on-chain funds to enter the liquidity marketplace in order to exchange their on-chain coins for off-chain coins. Once the trade has been atomically executed, the merchant immediately has usable inbound liquidity that can be used to accept payments from users. As the merchant acquires more liquidity and more channels in the future, they make additional contributions to the path diversity and strength of the network.

### 3.4 Cross-Chain Market Maker Liquidity Sourcing

As currency traders become more aware of the counterparty risk of trading on centralized exchanges, they become more motivated to find non-custodial exchange venues. The flexibility of channels on the Lightning Network make it a desirable platform for such a venue: channels allow for non-custodial trading at similar execution speeds to that of centralized exchanges. Additionally, channel based non-custodial exchanges are not vulnerable to front-running tactics executed by miners that can occur with other non-custodial concepts. Instead, the trade execution and even the prior trade history are *only* known to the participants, providing a greater degree of financial privacy.

Once again, we encounter a bootstrapping issue. How is a market maker on a payment channel-based non-custodial exchange meant to gather an initial pool of liquidity to service orders? We see CLMs as a natural solution. The market maker can seek out liquidity for relevant trading pairs by purchasing inbound channel liquidity, in addition to putting up its own outbound channel liquidity to other market makers. A balanced distribution of liquidity amongst market makers allows for new traders to participate in the exchange, knowing that their flows are balanced, meaning they can receive as much as they can send via the market maker, allowing them to instantly start to execute cross-chain atomic swaps.

### 3.5 Instant Lightning Wallet User On Boarding

Wallets commonly face the UX challenge of ensuring that a user can receive funds as soon as they set up a wallet. Some wallet providers have chosen to open new inbound channels to users themselves. This gives users the inbound bandwidth they need to receive, but can come at a high capital cost to the wallet provider as they need to commit funds at a 1:1 ratio. A CLM like Lightning Pool would allow wallet developers to lower their customer acquisition costs, as they would need to pay a relatively small amount relative to the volume of liquidity to be allocated to a new user. Just like the merchant purchasing inbound liquidity

in the above segment, a wallet provider could pay something like one thousand satoshis to have one million satoshis allocated to a user, instead of fronting the entire one million satoshis themselves.

### 3.6 Variance Reduction in Routing Node Revenue

Today, routing node operators aim to join the network in order to facilitate the transfer of payments as well as to earn fees over time by successfully facilitating payments. However, if a node isn't regularly routing payments (thereby earning a forwarding fee), then they aren't compensated for the various (though minor) risks they expose their capital to. With Pool, routing node operators are able to ensure that they're predictably compensated for the cost of their capital.

## 4 The Channel Lease Marketplace

In this section, we present an overview of the Channel Lease Marketplace architectural design. In section 6, we make a brief detour to define the *Shadowchain* application framework, before presenting a concrete instantiation of a CLM in the form of **Lightning Pool**.

### 4.1 High-Level Description

First, we describe our solution at a high-level. Drawing heavily from existing market auction design, we're interested specifically in double-call auctions that allow both the buyer and the seller to transact *indivisible* units of the good in question, which in this case is a channel lease. We then build upon this base double-call auction by leveling the informational playing field [6] by making all orders *sealed-bid*. Rather than the auction being cleared continually within a central-limit order book, we instead opt to utilize a discrete interval, frequent batched auction so as to mitigate front-running and other undesirable side effects [24]. Rather than participants paying what they bid (commonly called a pay-as-you-bid auction), all participants will instead pay the same *uniform clearing price*. Finally, all operations that result from a successful auction are batched and committed in a *single* atomic blockchain transaction.

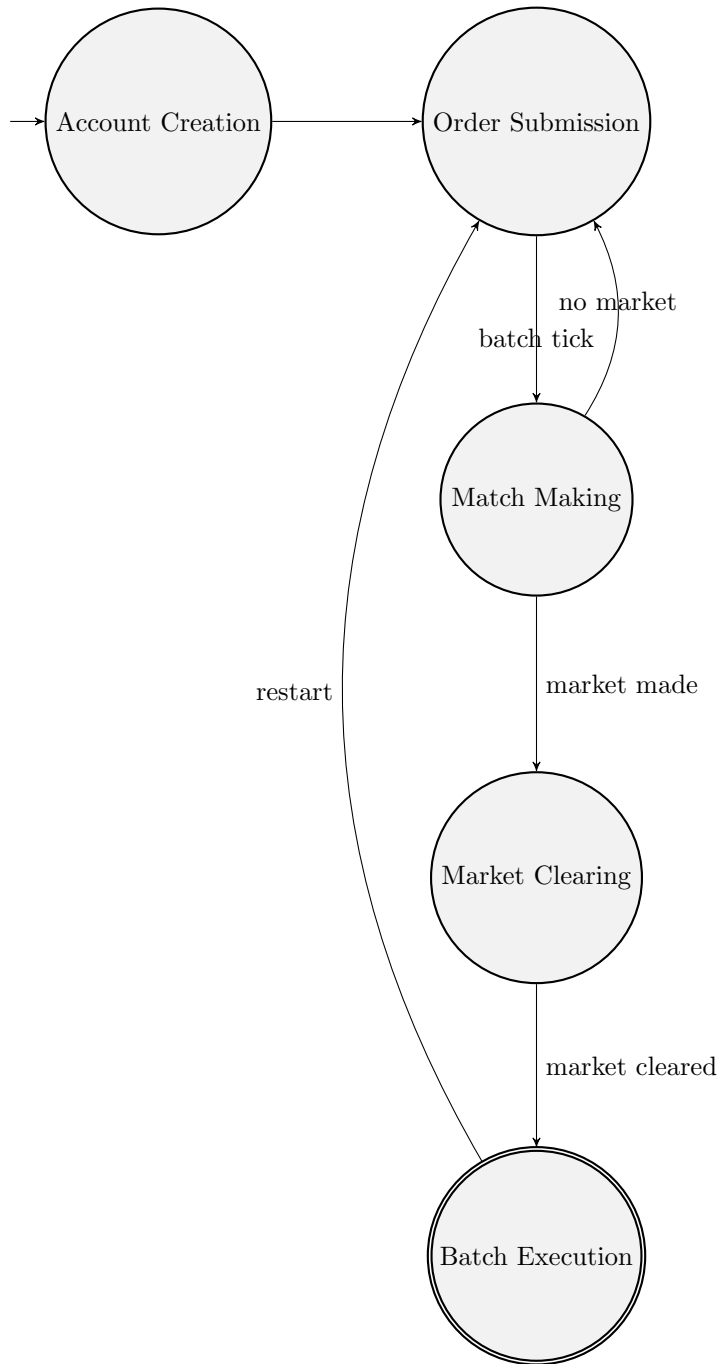


Figure 1: Auction State Machine

### Marketplace Auctioneer

We assume the existence of a *non-trusted* auctioneer  $\Lambda$  that publishes a master auctioneer key  $A_{\text{auction}}$  ahead of time. The auction itself is uniquely



identified by  $A_{\text{auction}}$  from the perspective of the system due to its *Shadowchain* qualities. The auctioneer implements a non-custodial auction via **Marketplace Accounts** that use a new unique key derived from  $A_{\text{auction}}$  as the second public key in the 2-of-2 multi-sig. The auctioneer accepts and validates orders off-chain, facilitates required account modifications, proposes a valid batch to each of the agents matched in an instance of the auction, and produces a batch execution transaction that creates the series of corresponding channel leases.

### Account Creation

Before being able to participate in the marketplace, we require that an agent first create a **Marketplace Account**. A **Marketplace Account** is a non-custodial account that forces an agent to commit capital in the form of Bitcoin to the market for a period of time. As we require agents to fully back all orders within an account, we eliminate a number of order spoofing vectors. Additionally, the time-locked, non-custodial nature of the account ensures a user is able to recover their funds fully without any additional on-chain transactions (aside from the sweeping transaction).

### Marketplace Order Units

We abstract over the base Bitcoin satoshi unit and define a **unit** from the point-of-view of the marketplace that serves as the base unit that all orders are expressed in and settled in. We assume that the value of a given **unit** is set such that even a single lease of the smallest unit is still economical from the perspective of the base blockchain and on-chain fees. All orders *must* be divisible by a whole unit, and the final clearing volume of a given batch is also expressed in these units.

### Order Submission

Once an agent has created a valid **Marketplace Account**, they can enter the order submission phase. It's important to note that this order submission takes place *off-chain*. Only the final execution of an auction batch takes place on-chain. During the order submission phase, agents are free to modify their accounts and orders. Only valid orders will be accepted as eligible for the next auction iteration.

### Market Clearing

Every  $\Upsilon$  minutes, the auctioneer attempts to *clear* the marketplace. An auction can be cleared if the lines of supply and demand cross such that at least a single unit is bought/sold. As the market has no explicit closing time, it's possible that during a market epoch, no market can be made. In the scenario that a market can be made, rather than each participant paying what they bid, the auctioneer instead uses a single clearing price based on the market's clearing price algorithm.

### Batch Execution

Once a market has been cleared, the batch execution phase begins. During this phase, the auctioneer sends a batch proposal  $\Pi$ , which describes the proposed market clearing structure.  $\Pi$  may either be a plaintext description of a valid clearing solution, or an "argument" describing one. Valid batches are then bundled into a single **Batch Execution Transaction** that updates all involved accounts, and creates any channel leases initiated within the batch. After a period of time  $\Upsilon$  has elapsed, the market is restarted with any new orders and accounts being considered for market clearing.

## 4.2 Lightning Channel Leases

### Liquidity Maker & Taker

We begin by introducing the concept of a **Liquidity Taker**:

**Definition 4.1.** (Liquidity Taker). A Liquidity Taker is an agent in a Channel Liquidity Market seeking to *obtain* new *inbound* channel liquidity of size  $A_{sat}$  for a period of  $T_{block}$  Bitcoin blocks.

A taker is prepared to either bootstrap the inbound liquidity with their own on-chain coins, or pay a *premium* in order to receive a "lease" of liquidity from another agent in the market. Takers populate the *demand* side of our market. They require new inbound liquidity in order to be able to immediately receive payments on the network, or to better position themselves as a routing node within the network.

A natural companion to the Liquidity Taker agent within a CLM is the **Liquidity Maker**:

**Definition 4.2.** (Liquidity Maker).

A Liquidity Maker is an agent in a Channel Liquidity Market seeking to earn *yield* by deploying up to  $A_{sat}$  Bitcoin into the Lightning Network for up to a period of  $T_{block}$  Bitcoin blocks, earning a profit  $\alpha$ .

Notice that we utilize Bitcoin block-time rather than wall-clock time (Median Past Time) [17] in these definitions, as we seek to enforce these durations using Bitcoin Script and using block-time is more objective compared to wall-clock time.

The profit ( $\alpha$ ) earned by a **Liquidity Maker** takes two forms:

- A one-time *premium*,  $R_{premium}$ , commanded by the Maker which reflects the latent demand and time-value of regular coins vs "lifted coins" (coins placed in channels).
- Ongoing recurring revenue,  $F_c$ , in the form of forwarding fees earned by facilitating payments to their matched taker.

We argue that the existence of such Channel Liquidity Markets will increase the *efficiency* of capital deployed to a payment channel network by allowing agents to signal the relative demand of lifted coins compared to non-lifted coins. Additionally, such markets also allow an existing routing node on the network to *re-allocate* lifted coins from a low-velocity section of the sub-graph, to one of higher velocity:

**Theorem 4.1** (Channel velocity revenue). Holding all channel liquidity equal, channels allocated to a higher velocity section of the sub-graph will yield a higher  $F_c$  than channel allocated to a low-velocity section of the sub-graph.

Intuitively, if each payment flow sourced at an incoming channel  $C_i$  and sunk at an outgoing channel  $C_o$  pays an equal forwarding fee per flow, then for a fixed unit of time, a higher velocity channel will result in higher total revenue in a time slice.

The role of Channel Liquidity Markets in a payment channel network is to reduce information asymmetry by allowing agents to signal their preferences for lifted coins vs non-lifted coins. The existence of *venues* where these markets can be carried out benefits the wider network by allowing agents to determine where their liquidity is most demanded on the network.

### Channel Leases

With our two primary agents defined, we now move on to the definition of a Lightning Channel Lease:

**Definition 4.3.** (Lightning Channel Lease). A Lightning Channel Lease is defined as,  $\Gamma = \{P_T, P_M, A_{sat}, D_{block}, r_i\}$ , where:

- $P_T$  is the `secp256k1` public key of the **Liquidity Taker**.
- $P_M$  the public key for the **Liquidity Maker**.
- $A_{sat}$  is the total amount of Bitcoin within the contract.
- $D_{block}$  is the duration of the contract expressed in Blocks.
- $r_i$  is the per-block interest rate as discovered in the  $i$ th instance of the market.

Note that the premium  $R_P$  as referenced above is parametrized in using the *lease duration*  $D_{blocks}$ :  $R_P(D_{blocks}) = r_i \cdot D_{blocks}$  as we deal in simple, rather than compounding, interest. The duration of the contract  $D_{blocks}$  is of great interest, as similar to U.S Treasury auctions, a *yield-curve* can be constructed based on the matched contents of a given auction iteration.

### 4.3 Non-Custodial Auction Accounts

In order to participate in the auction, we require all participants to deposit their trading balance into a Marketplace Account:

**Definition 4.4.** (Marketplace Account). A marketplace account is a *non-custodial* account defined as,  $\Psi = \{K_{sat}, T_{blocks}, P_{acct}, \Omega_{nodes}\}$  where.

- $K_{sat}$  is the total amount of Bitcoin available within the account.
- $T_{blocks}$  is the absolute expiry height of the account.
- $P_{acct}$  is a `secp256k1` public key that *uniquely* identifies the account.
- $\Omega_{nodes}$  is a set of Lightning Network nodes controlled by the account.

We stress that these accounts are *non-custodial* in that after a period of time  $T_{blocks}$  the agent is able to freely remove the funds from their account. Before this period has passed, an agent may require the participation of the auctioneer to close, deposit, or withdraw funds from their account. In the case of the **Liquidity Taker**, the funds within the account  $K_{sat}$ , must be sufficient to pay for any premium bids. Conversely for the **Liquidity Maker**, funds to be leased must be deposited into the account.

This structure, which forces all participants to fully commit all funds they wish to use within the marketplace into a non-custodial account, is similar to the concept of Fidelity Bonds [13]. This structure has a number of desirable properties including:

- **Order spoofing mitigation:** Within the CLM, as all orders must be "fully backed", it isn't possible to place a "fake" order that cannot be filled.
- **Time value opportunity cost:** By forcing all agents to suspend funds they wish to use within the market, those funds cannot be used elsewhere, thereby adding an implicit cost to joining the marketplace.
- **Deterministic batch execution construction:** As we'll see in later sections, the existence of a fixed account for each agent simplifies the clearing and execution process within the auction lifecycle.

These accounts in the abstract may take many forms, but as we focus on Bitcoin, as detailed in later sections, these accounts will take the form of a multi-signature output, with one key belonging to the auctioneer and the other belonging to the bidding participant.

#### 4.4 Order Structure & Verification

With our channel lease contract and account structure defined, we now move on to our order structure. As with any auction, orders are how the agents express their *preferences* with respect to what they wish to buy and sell. Importantly, all orders within the market must be backed by a valid non-expired account, and must carry an *authentication tag* which prevents order spoofing, and also ensures proper integrity of a given order once it has been submitted.

##### Order Structure

We define an **Order** within the context of a CLM as follows:

**Definition 4.5.** (Order). An **Order** is a authenticated  $n$ -tuple:  
 $\Theta = \{O_{type}, K_{nonce}, V_{ver}, P_{acct}, \Delta_{base}, \Delta_{aux}, T_{auth}\}$ , where:

- $O_{type} \in \{\text{Ask}, \text{Bid}\}$  denotes if an order is an Ask or a Bid. In addition to the version, this may affect how the  $\Delta_{aux}$  attribute is parsed.
- $K_{nonce}$  is an *order nonce* which uniquely identifies this order, and is typically derived as  $K_{nonce} = \leftarrow \$\mathbb{Z}_{2^\lambda}$ .

- $V_{ver}$  is the *version* of this order. As we'll see below, the version is used as an upgrade mechanism, and is needed in order to parse any newly added fields, as well as compute the digest required to check the authentication tag attached to an order.
- $P_{acct}$  is the public key that uniquely identifies this account.
- The set of base order details is:  
 $\Delta_{base} = \{\alpha_{rate}, A_{sat}, M_{pub}, L_{pub}, A_{addr}, C_{type}, D_{blocks}, F_{chain}\}$ , where:
  - $\alpha_{rate}$  is the desired *per-block* rate that the owner of the order wishes to buy or sell a channel lease at. Further below, this may be referred to as the BPY or block percentage yield.
  - $A_{sat}$  is the total contract size expressed in lease *units*. Restricting orders to whole units simplifies preference matching within the system.
  - $M_{pub}$  is the multi-sig public key to be used when creating the funding output of the channel.
  - $L_{pub}$  is the identity public key of the Lightning Node that wishes to buy/sell this channel.
  - $A_{addr}$  is the network address to be used to connect to the backing  $L_{pub}$  to initiate the channel funding process if this order is matched.
  - $C_{type}$  is the *type* of channel to be created if this order is matched.
  - $D_{blocks}$  is the target *lease duration* of the contract.
  - $F_{chain_{max}}$  is the max chain fee expressed in *sat/vbyte* that the owner of said order is willing to pay within a batch.
- The set of *auxiliary* details is implicitly defined by the order version  $V_{ver}$ .
- $T_{auth}$  is an authentication tag that allows the auctioneer, and other traders to validate the integrity and authenticity of the order.

An order allows a **Liquidity Taker** or a **Liquidity Maker** to express their *preference* with respect to what type of channel lease they're looking to buy or sell.

### Order Validation

Returning back to our tag  $T_{auth}$ , we will now specify how such a tag is to be computed, and verified. In the abstract, we require that the tagging scheme is **SUF-CMA** secure. Given this security requirement, we define two polynomial-time algorithms: (**GenOrderTag**, **VerifyOrderTag**) with the following requirements:

- **GenOrderTag**( $P_{acct_{priv}}, \Theta$ )  $\rightarrow T_{auth}$ . Given an input of the private key that corresponds to the public key of an account, and the complete order details, a valid tag  $T_{auth}$  is generated.
- **VerifyOrderTag**( $P_{acct}, \Theta, T_{auth}$ )  $\rightarrow b$ . Given a public key of an account holder, a valid tag, and the order itself, **VerifyOrderTag** outputs  $b = 1$  if the tag is valid.

As we use a public-key based tagging technique, the validity of an order is verifiable by any other active trader within the marketplace including the auctioneer of the marketplace.

## 4.5 Auction Design

In this section, we describe the abstract definition of a **Channel Liquidity Marketplace**, which addresses each of the issues presented in the bootstrapping section of the background, by creating a new form of batched auction which allows **Liquidity Takers** and **Liquidity Makers** to buy and sell Lightning Channel Leases in a non-custodial manner.

### 4.5.1 Auction Specification

We'll now specify the behavior and requirements of an using abstract **Channel Liquidity Marketplace** instance. We define the expected behavior and the client-facing interface of a **CLM** instance. A **CLM** is a tuple of polynomial-time algorithms divided into five distinct but related categories:

- System Initialization: **SystemInit**
- Account Operations: (**NewAccount**, **ModifyAccount**)
- Order Book Maintenance: (**SubmitOrder**, **CancelOrder**)
- Market Clearing: (**MatchMake**, **MarketClearingPrice**, **ClearMarket**)
- Batch Execution: (**ConstructBatch**, **ExecuteBatch**)

With behavior and semantics as expressed below.

#### System Initialization

Before the marketplace can be used, we require it to be initialized by the auctioneer. This initialization is a one-time process, and doesn't result in any trapdoor or "toxic waste" material being produced:

$\text{SystemInit}(1^\lambda, \Upsilon_{min}) \rightarrow (P_{\text{auction}_p}, P_{\text{auction}_s}, \Psi_A)$ . Denoting the security parameter as  $\lambda$ , the **SystemInit** algorithm takes as input the security parameter, and the batch interval  $\Upsilon_{min}$  expressed in minutes, and outputs a public ( $P_{\text{auction}_p}$ ) and private ( $P_{\text{auction}_s}$ ) key pair for the auctioneer. The auctioneer's public key will be used as a parameter in algorithms related to account creation, modification, and batch execution. This algorithm also returns  $\Psi_A$ , which is a special account owned by the auctioneer that will be used to collect fees, and during batch construction.

#### Account Operations

In order to create an account, agents will need to interact with the auctioneer itself. After account creation, an account can be modified freely (close, deposit, withdraw, etc.) if the account isn't part of an active batch:

$\text{NewAccount}(1^\lambda, P_{\text{auction}_p}) \rightarrow \Psi$ . The **NewAccount** algorithm takes as input our security parameter, and the auctioneer’s public key, and outputs a new account for the new agent within the marketplace. We require that all resulting accounts within the marketplace be *unique*. We permit a single logical agent to have multiple accounts.

$\text{ModifyAccount}(\Psi, P_{\text{auction}_p}) \rightarrow (b, \Psi')$ . The **ModifyAccount** algorithm takes an existing valid account  $\Psi$  and the auctioneer’s public key and performs an account modification. The algorithm returns  $b = 1$  if the modification was successful, and  $b = 0$  otherwise. An account modification may fail if the target account is already part of a pending batch. An account modification can either:

- Deposit new coins into the account.
- Withdraw coins from the account.
- Close the account by removing all coins from the account.

Note that as each of these operations requires an on-chain transaction, they can be freely batched with other on-chain transactions, or even the batch transaction that executes an auction.

### Order Book Maintenance

Once accounts in the marketplace are open, agents are able to submit orders between batch epochs. The *size* of all orders is expressed in units, and as we mention below, we permit partial matches of an order. A partial match can either update the order state in place, or require that the agent re-submit a new valid tag for the modified order in the batch execution phase:

$\text{SubmitOrder}(\Theta, T_{\text{auth}}) \rightarrow b$ . The **SubmitOrder** algorithm takes as input a structurally sound order  $\Theta$ , and its authentication tag  $T_{\text{auth}}$  and outputs a bit  $b$ . The bit  $b = 1$  if the order is valid according to market place rules, and the **VerifyOrderTag** returns  $b = 1$  given the specified parameters.

$\text{CancelOrder}(\Theta, K'_{\text{nonce}}) \rightarrow b$ . The **CancelOrder** given an existing order  $\Theta$  and the *opening* of the  $K_{\text{nonce}}$  commitment  $K'_{\text{nonce}}$  and returns  $b = 1$  if the commitment opening is valid, and there exists an order identified by the base  $K_{\text{nonce}}$  value.

### Market Clearing

Once all orders have been placed and the batch interval of  $\Upsilon$  has elapsed, the auctioneer will attempt to clear the market using the following algorithms:

$\text{MatchMake}(\{\Theta_0, \dots, \Theta_n\}) \rightarrow \Phi_b = \{(\Theta_{b_0}, \Theta_{a_0}), \dots, (\Theta_{b_n}, \Theta_{a_n})\}$ . The **MatchMake** algorithm takes as input the set of valid orders submitted during the past batch interval and outputs a series of tuples that reflect properly matched orders.  $\Theta_a$  represents an order with  $O_{\text{type}} = \text{Ask}$ , while  $\Theta_b$  represents an order with  $O_{\text{type}} = \text{Bid}$ . Note that since we allow *partial* matches, a given order may appear multiple times in the final match set. We require that a valid implementation be able to perform proper *multi-attribute* matching due to the existence

of the  $\Delta_{aux}$  portion of an order’s structure.

**MarketClearingPrice**( $\Phi_b$ )  $\rightarrow c_{price}$ . The **MarketClearingPrice** algorithm accepts the set of orders matched by the **MatchMake** algorithm and returns the *market clearing price* of the prior batch. The precise market clearing price algorithm is left as a free parameter, with algorithms such as first-rejected-bid or last-accepted-bid likely being used. Utilizing a single market clearing price is intended to promote fairness and has a number of additional benefits over pay-as-bid auctions [23].

**ClearMarket**( $\Psi_A, \Phi_b, \{\Psi_0, \dots, \Psi_n\}, c_{price}$ )  $\rightarrow (\Psi'_A, \{\Gamma_0, \dots, \Gamma_n\}, \{\Psi'_0, \dots, \Psi'_n\})$ . The **ClearMarket** algorithm takes as input a prior set of matched orders within a batch, the auctioneer’s account, the set of accounts involved in the batch, and the market clearing price of a given batch. The algorithm outputs a set of channel leases to be created by a batch along, a set of updated accounts, and an updated version of the auctioneer’s account with any trading fees accrued during market clearing.

- As shorthand, we use  $\Delta_i$  to refer to a cleared batch (the set of resulting accounts after the updates have been made to produce the set of desired channel leases).

### Batch Execution

Once we’ve been able to make a market, and have the description of the resulting market state (the accounts and the channel leases to be created), we can now move on to *executing* the resulting batch. We use the following algorithms to do so:

**ConstructBatch**( $\Delta_i$ )  $\rightarrow B_{t_i}$ . The **ConstructBatch** algorithm takes a valid market clearing, which can be seen as a delta on the auction state, and returns a valid transaction that *atomically* executes the given batch on the blockchain.

**ExecuteBatch**( $B_{t_i}$ )  $\rightarrow b =$ . The **ExecuteBatch** algorithm takes a fully valid batch and attempts to commit it by confirming the transaction in the target base blockchain. Once the batch has been confirmed, all operations contained within a batch are considered executed, and can be used as inputs to additional iterations of the auction life cycle.

## 5 The Shadowchain: A Bitcoin Overlay Application Framework

In this section, we present the concept of a **Shadowchain**, a non-custodial application overlay framework that we’ll use to construct a concrete instantiation of a CLM. We note that shadowchains may also be of independent interest, as they’re a novel way to layer more complex interactions on top of the base Bitcoin blockchain. **Shadowchains as we present them can be implemented on the base Bitcoin blockchain today without any additional changes or enhancements.**



However, further extensions to Bitcoin such as cross-input signature aggregation and covenants could serve to dramatically improve scalability properties of shadowchains.

## 5.1 High-Level Description

First, we provide a high-level description of the shadowchain application framework.

**The Shadowchain Usecase.** A shadowchain can be used to implement non-custodial smart contract systems on top of the base Bitcoin blockchain. Typically, one would opt for a shadowchain if the complexity of the state transition logic of the smart contract system could not be fully expressed using the base Bitcoin Script. Shadowchains allow an application designer to use the Bitcoin blockchain for censorship resistant settlement, while pushing the more complex portions of the application (state, logic, etc.) *off-chain*.

**Shadowchain Roles & Lifted UTXOs.** A shadowchain has two primary classes of agents: users, and the orchestrator. The orchestrator defines the state transition function of the shadowchain, a set of non-trusted initialization parameters, and upgrade mechanisms. A user is able to join a shadowchain by "lifting" their UTXOs into the higher-level shadowchain. The process of lifting (defined further below) entails the user placing funds within a time-lock released, multi-signature output that enforces cooperation between the user and the shadowchain orchestrator.

**Shadowchain Operation.** The shadowchain orchestrator accepts transaction data from users and periodically proposes a new *shadowchain block*. A shadowchain block takes as input the set of Lifted UTXOs that accepted the latest block proposals, and produces a set of new UTXOs, which are the end state after the state transition function has been evaluated. A shadowchain is even permitted to use *multiple* distinct state transition functions. As user funds cannot move without both multi-sig signatures, users are able to fully validate (possibly using techniques such as zero knowledge proofs that the resulting UTXO state was properly derived from the known state transition function. Note that due to this structure, complex "exit-games", or fraud proofs are not required as a user simply won't sign off on a fraudulent state, and a user's UTXO is always manifested (in its base form) on the main blockchain.

**Ephemeral Lifted UTXOs.** In the scenario that the shadowchain orchestrator disappears, or is unresponsive, users are able convert their lifted UTXO into regular ones by spending their coins after the time-lock has expired. This construct of an *ephemeral* lifted UTXO has a number of desirable properties on the application level, as the time-locked commitment of funds can serve to mitigate a number of application-level issues such as spam or sybil resistance.

**Shadowchain Cut-Through** As the evolution of a state transition function happens *off-chain*, it's possible to coalesce several distinct shadowchain blocks into a single block that combines successive invocations of the state transition function. This technique is similar to transaction cut-through [18] but is performed in a multi-party setting. Leveraging this technique, the shadowchain orchestrator can optimistically treat the current latest shadowchain transaction in the mempool as an in-memory data structure to be updated off-chain (via transaction replacement techniques), with the state being "written to disk"

once confirmed. As a result, it's possible to commit several shadowchain states (possibly hundreds) in a single logical Bitcoin transaction.

**Shadowchain Upgrades.** Finally, similar to the base blockchain, a shadowchain can also be *upgraded* in a forward and backward compatible manner. In other words, it's possible for a shadowchain orchestrator to *soft-fork* the state transition logic by restricting a valid state transition to enable new behavior. Notably, the orchestrator can do this in a desynchronized manner such that only those wishing to use the features of the new state transition function need to adhere to the new rules. Additionally, an orchestrator can opt introduce new backward incompatible state transition functions. Note that because of the batching capabilities inherent in Bitcoin transactions, an orchestrator can commit multiple logical shadowchain blocks (with distinct state transition functions) in a single atomic Bitcoin transaction.

To summarize, the shadowchain application framework is a novel technique for constructing overlay applications on the base Bitcoin blockchain in a non-custodial manner. Shadowchains avoid the complexity of fraud proofs and exit games by ensuring that the user has custody of their funds at all times and is able to fully validate any proposed state transition. Shadowchains are able to compress several logical state transitions into a single Bitcoin transaction using a multi-party cut-through technique. An orchestrator of a shadowchain is also able to upgrade the state transition logic on the fly, in a purely off-chain manner.

## 5.2 Comparison To Related Frameworks

### 5.3 The Shadowchain Framework

In this section, we present the abstract shadowchain application framework. Applications are intended to use this framework, providing implementations of specified virtual functions to fully specify and execute their application.

#### 5.3.1 Shadowchain Orchestrator

First, we introduce the glue that keeps a shadowchain together, the orchestrator:

**Definition 5.1.** (Orchestrator). The **Orchestrator** is a non-trusted entity at the root of a shadowchain, parametrized by its long-term public key:  $O_{chain} = P_O$ . The duty of an **Orchestrator** is to propose new blocks (the result of a state transition) to the set of live **Lifted** UTXOs that make up the shadowchain.

A given **Orchestrator** is a non-trusted entity, and can be uniquely identified by its long-term public key. The long-term public key  $P_O$  can also be used to uniquely identify a given shadowchain, similar to the Genesis Block hash of a normal blockchain.

#### 5.3.2 Lifted UTXOs

Next, we define the **Lifted** UTXO, which is the representation of a user's state within a given shadow chain:

**Definition 5.2.** (Lifted UTXO). A **Lifted** UTXO is a tuple,  $\phi_U = (A_{sat}, T_{expiry}, P_u, P_o)$ , where:

- $A_{sat}$  is the size of a **Lifted UTXO** expressed in *satoshis*.
- $T_{expiry}$  is the *absolute* expiry height of the **L0**, after which the owner is able to unilaterally move the funds back to the "base" Bitcoin blockchain.
- $P_u$  is the public key of the end user, which is 1/2 of the public keys used in the public key script of the output which manifests this **L0** on the base blockchain.
- $P_o$  is the public key of the **Orchestrator**, typically derived from its base long-term key  $P_O$ . This key will be used as the other half of the multi-sig script of the on-chain manifestation of the **L0**.

The construct of a **Lifted UTXO** is similar to the existing concept of a Fidelity Bond [13] with an application-specific twist. This process is akin to creating a new 'account' within a Shadowchain. The time-lock release nature of the UTXO means that a user can always recover funds if the **Orchestrator** becomes unresponsive. In addition to this, a natural cost in the form of chain fees is added, which increases the barrier for potentially malicious users to interact with the shadow chain.

### 5.3.3 The Shadowchain

In this section, we present the abstract definition of a shadowchain, building upon the definition provided above. In addition to this, we describe the typical shadowchain lifecycle using the aid of some additional helper functions, which are also intended to be included in the application logic of the shadowchain.

#### Shadowchain Components

First, we define the core components of the shadowchain.

**Definition 5.3.** (Shadowchain). An instantiation of a **Shadowchain** is defined as a tuple:  $\Sigma = (U_L, U_O, \Delta_F, E_{exe}, A_T)$ , where:

- $U_L = \{\phi_i, \dots, \phi_n\}$  is the set of non-expired **Lifted UTXOs** observed by the orchestrator.
- $U_0$  is the current UTXO of the orchestrator, where they may accrue application level fees.
- $\Delta_F = \{\Delta_{f_0}, \dots, \Delta_{f_n}\}$  is the set of current state transition functions.
- $E_{exe}$  is the abstract execution environment of the shadowchain which all participants will use to verify the correctness of a proposed state transition.
- $A_T$  is the abstract form of the structure of the higher-level application's fundamental transaction.

#### Shadowchain Algorithms

Given the above components, we define the operation of a shadowchain using a series of polynomial-time algorithms segmented into the following logical categories:

- System Initialization: **InitChain**
- UTXO Management: (**LiftUTXO**, **UnliftUTXO**, **ExitChain**)
- Block Proposal & Validation: (**ConstructBlock**, **ProposeBlock**)
- Chain Execution: **CommitBlock**
- Block Cut-Through: **CoalesceBlocks**
- Chain Upgrade: **UpgradeChain**

With behavior and semantics as expressed below.

### System Initialization

Before a shadowchain can be used for a given application, the system must be initialized. This process results in the creation of the orchestrator's long-term public key, the execution environment, and the set of state transition functions:

**InitChain**( $1^\lambda$ )  $\rightarrow (U_0, P_0, \Delta_F, E_{exe})$ . Given the security parameter  $\lambda$  (expressed in unary), the **InitChain** method returns the initial self-lifted UTXO of the orchestrator, the long-term public key of the orchestrator, and the set of initial state transition functions along with the starting execution environment.

### UTXO Management

Once a shadowchain has been initialized with a given set of parameters, users can begin lifting their UTXOs, enabling participation in shadowchain blocks and operations. The process of entering the shadowchain is referred to as UTXO Lifting, while exiting is the reverse process:

**LiftUTXO**( $T_{expiry}, \{U_{N_0}, \dots, U_{N_n}\}, P_0$ )  $\rightarrow \phi_U$ . The **LiftUTXO** algorithm takes a series of normal UTXOs ( $U_N$ ), the absolute expiration height of the UTXO, and the long-term public key of the orchestrator  $P_0$ , outputting a new **Lifted UTXO** for the target user. The total value of the set of input UTXOs must be greater-than-or-equal to the value of the resulting **Lifted UTXO**.

**ExitChain**( $\phi_U, B_{height}$ )  $\rightarrow U_N$ . Given a lifted UTXO with expiration height  $T_{expiry} > B_{height}$ , where  $B_{height}$ , the **ExitChain** method spends an existing lifted UTXO and resolves the user's funds in a regular unencumbered UTXO. Users will use this algorithm if they wish to exit the chain in cases where the orchestrator is no longer being responsive.

**UnliftUTXO**( $\phi_U$ )  $\rightarrow U_N$ . Given a lifted UTXO, the **UnliftUTXO** method create a new normal un-lifted UTXO that returns all funds to the user. This is the optimistic version of the **ExitChain** algorithm, in that it requires cooperation from the orchestrator as the time-release clause of the Lifted UTXO's script is not yet unlocked.

### Block Proposal & Validation

Once a shadowchain has a sufficient number of lifted UTXOs and the system has been fully initialized, block proposal and validation can commence. This process is similar to the process of nodes broadcasting transactions, and miners ordering them within blocks in the base Bitcoin system:

**ConstructBlock**( $\phi_{live}, T_{xn}, E_{exe}, \Delta_F$ )  $\rightarrow B_S$ . Given inputs of the set of 'live' Lifted UTXOs, the set of transactions belonging to the live UTXOs, the execution environment, and the current set of valid state transition functions, the **ConstructBlock** outputs a valid shadowchain block to extend the main chain where:

- $\phi_{live} = live(\{\phi_{U_0}, \dots, \phi_{U_N}\})$  is the set of 'live' Lifted UTXOs where the *live* algorithm uses a heartbeat-like protocol to detect the current set of active users.
- $T_{xn}$  is the application-specific transaction format used within the shadowchain.
- $B_S = (T_{xn}, \{\phi_{U_0}, \dots, \phi_{U_N}\}, \Delta_f, \{\phi'_{U_0}, \dots, \phi'_{U_N}\}, U_A)$ , is the shadowchain block itself, which is composed of the set of application transaction, input Lifted UTXOs, the resulting output UTXOs after applying the set of state transition functions, and  $U_A$  any new application-specific UTXOs produced as a result of the state transition function. Lifted UTXOs can be consumed fully by the state transitions, therefore  $|\{\phi_{U_0}, \dots, \phi_{U_N}\}| \geq |\{\phi'_{U_0}, \dots, \phi'_{U_N}\}|$  must be given.

Once a block has been constructed, the orchestrator of the shadowchain now must propose said block to the set of live Lifted UTXOs before it can move onto the next phase of shadowchain operation. As a given user may reject a block, either implicitly due to being offline, or explicitly due to a violation of the shadowchain consensus rules, this phase may be repeated a number of times. The operator will use the following algorithm to propose blocks:

**ProposeBlock**( $B_S, \phi_{live}$ )  $\rightarrow b$ . The **ProposeBlock** attempts to propose the given shadowchain block to the set of live Lifted UTXOs. The algorithm returns  $b = 1$  if all of the participants accept the block. Once all participants have accepted the block, we can now proceed to the execution and block commitment phase.

### Chain Execution

Once the operator has established a stable set of participants that accept the proposed shadowchain block, it can execute the block and commit it in the base Bitcoin blockchain:

**CommitBlock**( $B_S$ )  $\rightarrow (b, TX_{id})$ . The **CommitBlock** takes a valid shadowchain block, and attempts to obtain all necessary witnesses to unlock the Lifted UTXOs to re-create them post state function application along with any new application-specific UTXOs. The algorithm returns  $b = 1$  if the operator was able to successfully obtain all necessary witnesses, and broadcasts the Bitcoin transaction that commits the shadowchain block. Otherwise, the operator may need to re-propose a new block to a subset of the live Lifted UTXOs.

## Block Cut-Through

Given the structure of shadowchain blocks and state transition functions, it's possible for a shadowchain orchestrator to compress several distinct shadowchain blocks into a single instance that is equivalent to the application of successive state transition functions on a series of distinct shadowchain blocks:

$\text{CoalesceBlocks}(\{B_{S_0}, \dots, B_{S_N}\}) \rightarrow B'_S$ . Given a series of *consecutive* shadowchain blocks, the **CoalesceBlocks** algorithm compresses each consecutive block into a single block  $B_{S_N}$  that has an equivalent output state to the serial application of the state transition functions on each individual shadowchain block.

The **CoalesceBlocks** algorithm is similar to the concept of transaction cut-through [18] for UTXO-based blockchains. Using this algorithm, the operator is able to propose a new equivalent block which produces the same set of Lifted UTXO outputs along with any other application-specific outputs produced by any of the intermediate blocks. This can be done post-facto, and also in an optimistic manner in order to coalesce several unconfirmed shadowchain blocks into a single one.

## Chain Upgrade

The process of upgrading the chain in terms of the *types* of application-level transactions offered, or the set of valid state transition functions used, can be done entirely off-chain in an asynchronous manner. In order to update the environment, state transition functions, or the application-level transactions, the orchestrator simply needs to utilize the following algorithm:

$\text{UpgradeChain}(\Delta_{new}, E'_{exe}, T'_{xn}) \rightarrow \perp$ . The **UpgradeChain** is an algorithm that's executed entirely in an off-chain manner allowing the operator of the shadowchain to upgrade some or all of: the application transaction data, the execution environment, and the set of state transition functions. Due to the nature of the shadowchain, users of this new functionality each user must update their local state. However, note that the Lifted UTXO remains the same, as the operator's long-term public key remains unchanged.

### 5.3.4 Shadowchain Operation

In this section, we'll put together the above algorithms to outline the main execution loop of the shadowchain from the perspective of the orchestrator as well as the participants. Shadowchain operation can be viewed as a linear deterministic state machine that uses the main Bitcoin blockchain to transition between states.

## Orchestrator State Machine

The main Orchestrator State Machine loop first attempts to process any new UTXO lifting requests and will optimistically attempt to merge any existing unconfirmed shadowchain blocks that can be coalesced. Independent of either of these clauses, the system will then enter into its main loop from which

the orchestrator will attempt to construct a new block. This loop will iterate through the set of live transactions, propose a block to the set of live UTXOs, filter any participants that reject the block, then attempt to commit the new block to the blockchain.

```

OrchestrateChain( $\Delta_F, E_{exe}$ )
1: repeat
2:   if haveNewLiftReqs()
3:      $\{\phi_{\text{new}}\} \leftarrow \text{liftNewUTXOs}(P_O)$ 
4:      $\{\phi_U\} \leftarrow \phi_U \cup \phi_{\text{new}}$ 
5:   if numUnconfBlocks() > 1
6:      $B'_S \leftarrow \text{CoalesceBlocks}(\text{unconfBlocks}())$ 
7:      $(b, txid) \leftarrow \text{CommitBlock}(B'_S)$ 
8:      $T_{xn} \leftarrow \text{liveTransactions}()$ 
9:      $\phi_{\text{live}} \leftarrow \text{liveLiftedUTxos}()$ 
10:     $b \leftarrow 0$ 
11:    while  $b == 0 \& \text{len}(\phi_{\text{live}}) > 0$ 
12:       $B_S \leftarrow \text{ConstructBlock}(\phi_{\text{live}}, T_{xn}, E_{exe}, \Delta_F)$ 
13:       $b \leftarrow \text{ProposeBlock}(\phi_{\text{live}}, B_S)$ 
14:      if  $b == 0$ 
15:         $\phi_{\text{live}} \leftarrow \text{filterRejects}(\phi_{\text{live}})$ 
16:      if  $b == 1$ 
17:         $(b', TX_{id}) \leftarrow \text{CommitBlock}(B_S)$ 
18:

```

### Participant State Machine

The main state machine of each shadowchain participant is essentially a mirror of the orchestrator state machine. First, it will process any requests to modify its existing Lifted UTXO, gather any unconfirmed transactions, then await a new block proposal from the orchestrator.

```

ExtendChain( $\phi_U, \Delta_F, E_{exe}$ )
1: repeat
2:    $T'_{xn} \leftarrow \text{unconfTxns}()$ 
3:   submitTxns( $T_{xn}$ )
4:    $B'_S \leftarrow \text{await } \text{recvBlockProposal}()$ 
5:    $b \leftarrow \text{ValidateBlock}(B'_S)$ 
6:   if  $b == 1$ 
7:     sendWitnesses( $U_O$ )
8:      $b' \leftarrow \text{await } \text{blockFinalize}$ 
9:     if  $b' == 1$ 
10:       localCommitBlock()

```

## 6 Lightning Pool: A Channel Liquidity Marketplace as a Shadow Chain

In this section, we build upon the prior sections outlining the abstract **Channel Lease Marketplace** definition, as well as shadowchain operation, and construct our **Lightning Pool** implementation at a low-level. We first begin by detailing our implementation of the CLM algorithms, as well as our choice of certain free parameters. With this concrete structure in place, we'll then go up a layer of abstraction to demonstrate how Lightning Pool can be operated as a shadowchain on Bitcoin today, without any further modifications enhancements.

### 6.1 Instantiating a CLM

#### 6.1.1 System Initialization

##### Batch Key Parameter

Before an instance of Lightning Pool can be used by willing agents, the system must first be initialized. This operation can be performed only by the Orchestrator of the auctioneer. Within the system, we'll utilize an incremented Elliptic Curve point which we refer to as the **batchID** for several operations. The **batchID** serves to uniquely identified a given batch, and is incremented after each successful batch.

The **batchID** itself is a nothing up my sleeve (NUMS) point which has been generated in a manner that no one, not even the auctioneer knows the discrete log to. The raw serialized **batchID** (for the very first batch) within the **Lightning Pool** system can be expressed in the following syntax (displayed using hex-encoding of a compressed key encoding of the batch key itself):

$$B_{k_0} = NUMS_{gen}(1^\lambda)$$

Let the current **batchID** for the  $n^{\text{th}}$  batch be  $B_{k_i}$ . Let  $\mathbb{G}$  be a cyclic group of order prime order  $p$ , generated by an element  $G$ . We now define two helper functions to allow us to "seek" round the **batchID** key-space:

<u>IncrementBatchKey(<math>B_{k_i}</math>)</u>	<u>DecrementBatchKey(<math>B_{k_i}</math>)</u>
1 : <b>return</b> $B_{k_i} + G$	1 : <b>return</b> $B_{k_i} - G$
<u>InitBatchKey</u> 1 : <b>return</b> $B_{k_0}$	

With the system initialization , we'll now move onto specifying the structure of a **Marketplace Account** within **Lightning Pool**.

##### Auctioneer State Initialization



Now that we have defined the set of initialization and manipulation methods for our batch key parameter, we'll move onto the initialization of the remaining system. In order to prevent key-reuse across the entire system, we employ a similar key-derivation scheme to that of the Lightning Network's current commitment transaction format [21]. This key derivation will be used within all account scripts within the system, as well as the auctioneer's main account.

First, we define a helper function for deriving the auctioneer's current key  $A_{pk_i}$  from the static auctioneer account  $keyA_{pk}$  parameterized by the current **batchID**. As noted above, the **batchID** serves as both a public key within the system as well as a counter. The **batchID** may be expressed as a normal compressed public key, or as an integer  $N_b$  which denotes the scalar multiple off-set from the starting batch key  $B_{k_0}$ :  $[B_{k_0}]N_b$ . We define the **auctioneerScript** as follows:

$\text{auctioneerScript}(A_{pk}, B_{k_i})$ <hr style="border: 0; border-top: 1px solid black; margin: 5px 0;"/> 1 : $A_{pk_i} \leftarrow A_{pk} + G \cdot H(A_{pk} \  B_{k_i})$ 2 : <b>return</b> OP_CHECKSIG $A_{pk_i}$
---

The script itself is a simple script that simply verifies a proper auctioneer signature given a particular **batchID**. In order to make all scripts uniform (as the account scripts are P2WSH outputs, we wrap this script in a P2WSH layer of indirection as well. With this algorithm defined, we can now define the **SystemInit** implementation for **Lightning Pool** that derives the first batch key along with the starting script of the auctioneer.

$\text{SystemInit}(1^\lambda, \Upsilon_{min})$ <hr style="border: 0; border-top: 1px solid black; margin: 5px 0;"/> 1 : $batchKey \leftarrow \text{InitBatchKey}()$ 2 : $A_{sk} \leftarrow \$\mathbb{Z}_{2^\lambda}$ 3 : $A_{pk} \leftarrow GA_{sk}$ 4 : $aScript \leftarrow \text{auctioneerScript}(A_{pk}, batchKey)$ 5 : $pkScript \leftarrow \text{p2wsh}(aScript)$ 6 : <b>return</b> $(A_{sk}, A_{pk}, \text{utxo}(pkScript, seedSats))$
--

Once the third return value, auctioneer's master account has been confirmed. Participants are able to open accounts, submit orders, and finally participate in auction batches.

### 6.1.2 Lightning Pool Accounts

Next, we move onto the **Marketplace Account** structure for agents within the **Lightning Pool** marketplace itself. Similar to the auctioneer's master account, we apply a key derivation scheme that combines the auctioneer's key, the batch key, a key supplied by the trader, and a distinct trader specific secret to ensure, that:

- A P2WSH output script is *never* re-used.
- Traders rotate keys with each batch.

- The set of trader keys using within a batch itself is indistinguishable w.r.t the input keys referenced and newly created outputs keys.

Let  $P_t$  be a trader's base key,  $S_t \leftarrow \mathbb{Z}_2^\lambda$  be an account-specific secret generated by the trader, and  $P_A$  be the auctioneer's long-term public key. We define a new algorithm **traderAccountScript** which will be used to derive the **pkScript** for a given trader:

<pre> <b>traderAcctScript</b>(<math>P_t, S_t, P_A, B_{k_i}</math>) 1 : <math>t \leftarrow H(B_{k_i}    S_t    P_t)</math> 2 : <math>P'_t \leftarrow P_t + (t \cdot G)</math> 3 : <math>P'_A \leftarrow P_A + H(P'_t    P_A) \cdot G</math> 4 : <b>witnessScript</b> <math>\leftarrow \{</math> 5 :   <math>P'_t</math> <b>OP_CHECKSIGVERIFY</b> 6 :   <math>P'_A</math> <b>OP_CHECKSIG</b> 7 :   <b>OP_IFDUP</b> 8 :   <b>OP_NOTIF</b> 9 :   <math>T_{blocks}</math> <b>OP_CHECKLOCKTIMEVERIFY</b> 10 :  <b>OP_ENDIF</b> 11 : <math>\}</math> 12 : <b>return witnessScript</b> </pre>
---

The above script can either be spent via the time out clause using the following witness stack, using **nil** value passes an empty signature to force execution of the timeout clause.

**nil traderSig witnessScript**

Or via the normal spend path way which will be used to authorized batches, account closing, and any other account modifications:

**auctioneerSig traderSig witnessScript**

Note that each trader starts using the current batch key at the time they joined the marketplace. However, as we'll see below in the execution section, a trader's key gets rotated with each batch they participate in, meaning that the set of batch keys used within a trader's account output script will eventually become de-synchronized across the market unless all trader's participate in all batches, which is unlikely.

Given the **traderAccountScript** algorithm, we'll now define our implementation of the set of **Account Operations** methods:

<b>NewAccount</b> ( $1^\lambda, P_{\text{auction}_p}$ )
<pre> 1: <math>k_t \leftarrow \\$\mathbb{Z}_{2^\lambda}</math> 2: <math>P_t \leftarrow k_t \cdot G</math> 3: <math>\text{acctScript} \leftarrow \text{traderAcctScript}(P_t, S_t, P_A, B_{k_i})</math> 4: <math>\text{accountTxn} \leftarrow \text{makeTxn}(\text{acctScript})</math> 5: <math>\text{traderSignTxn}(\text{accountTxn})</math> 6: <math>\text{broadcastTxn}(\text{accountTxn})</math> 7: <math>\Psi \leftarrow \text{await confirmation}(\text{accountTxn})</math> 8: <b>return</b> <math>\Psi</math> </pre>
<b>ModifyAccount</b> ( $\Psi, P_{\text{auction}_p}$ )
<pre> 1: <b>if</b> <math>\text{inPendingBatch}(\Psi)</math> 2:   <b>return</b> <math>(0, \Psi)</math> 3: 4: <math>B_{k_i+1} \leftarrow \text{IncrementBatchKey}(\Psi, B_{k_i})</math> 5: <math>\text{acctScript} \leftarrow \text{traderAcctScript}(P_t, S_t, P_A, B_{k_i+1})</math> 6: 7: <math>\text{txInputs} \leftarrow \text{readInput clientStream}</math> 8: <math>\text{txOutputs} \leftarrow \text{readInput clientStream}</math> 9: <math>\text{action} = \perp</math> 10: <b>match</b> : 11:   <math>\text{acctScript not in txOutputs}</math> : 12:     <math>\text{action} = \text{CLOSE}</math> 13:   <math>\text{txOutputs.acctOutput.Value} &gt; \Psi.\text{Value}</math> : 14:     <math>\text{action} = \text{DEPOSIT}</math> 15:   <math>\text{txOutputs.acctOutput.Value} &lt; \Psi.\text{Value}</math> : 16:     <math>\text{action} = \text{WITHDRAW}</math> 17: 18: <math>\text{accountTxn} \leftarrow \text{makeTxn}(\text{acctScript}, \text{action}, \text{txInputs}, \text{txOutputs})</math> 19: <math>\text{traderSignTxn}(\text{accountTxn})</math> 20: <math>\text{broadcastTxn}(\text{accountTxn})</math> 21: 22: <math>\Psi' \leftarrow \text{await confirmation}(\text{accountTxn})</math> 23: <b>return</b> <math>\Psi'</math> </pre>

Notice how the `batchKey` is incremented for all account modification operations. In this manner, the `batchKey` also serves as a sequence number within the scripts to ensure that no scripts are re-used within the system.

Users are able to recover their accounts themselves if data has been lost by scanning the chain for the above instances of `newAcctScript` and utilizing the timeout clause within the account script of a trader's account.

### 6.1.3 Channel Leases in the Lightning Network

Now that we have our concrete account structure, we'll move on to presenting a concrete instantiation of a Channel Lease based on today's widely used payment channel implementation within Lightning Network.

#### Channel Lease Duration Enforcement

The unique component that sets apart a regular channel from one that was created via a channel lease contract is *duration* enforcement. As a channel lease contract states the capital must be committed to the network for a minimum period of time, in order to implement this in a trust-minimized manner, we must enhance the existing channel format [21] used in the Lightning Network today. Our tool of choice for creating minimum duration enforced channels is the `OP_CHECKLOCKTIMEVERIFY` [29] op-code. Our construction is similar to the one proposed in [30] Minimally, we need to enforce the following conditions:

- The **Liquidity Maker** involved in a channel lease cannot sweep the funds in their settled commitment output as manifested on their commitment or the commitment of the remote party until  $D_{block}$  Bitcoin blocks has passed since the creation of the channel lease.
- The **Liquidity Maker** also cannot fully sweep any funds that are suspended within HTLC outputs until  $D_{block}$  Bitcoin blocks has passed.

The second item is of great importance to ensure that the seller of a channel lease can't just move all the committed funds into HTLCs, then close the channel and be fully refunded, netting the lease premium in the process. Instead, we need to ensure that the node is able to resolve any HTLCs on-chain if needed, while still being forced to commit the funds in ancestors of the multi-sig funding output, until the lease duration has expired.

We now make a small modification first to the settled local output of the **Liquidity Maker**:

```
OP_IF
  <revoke key>
OP_ELSE
  <lease duration in blocks>
  OP_CHECKLOCKTIMEVERIFY
  OP_DROP

  <delay in blocks>
  OP_CHECKSEQUENCEVERIFY
  OP_DROP

  <delay key>
OP_ENDIF

OP_CHECKSIG
```

Next, we make a similar modification to the settled remote output of the **Liquidity Maker**, assuming anchor output based channels [4] are used:

```
<localKey> OP_CHECKSIGVERIFY
<lease duration> OP_CHECKLOCKTIMEVERIFY
1 OP_CHECKSEQUENCEVERIFY
```

Finally, we modify the offered HTLC outputs of the **Liquidity Maker** for their local commitment transaction:

```
OP_IF
  <revoke key>
OP_ELSE
  <lease duration in blocks>
  OP_CHECKLOCKTIMEVERIFY
  OP_DROP

  <delay in blocks>
  OP_CHECKSEQUENCEVERIFY
  OP_DROP

  <delay key>
OP_ENDIF

OP_CHECKSIG
```

A less trust-minimized version of channel lease duration enforcement is possible simply by having the **Liquidity Taker** refuse a cooperative channel closure until the lease duration has expired. With this non-script based enforcement, the only direct option the **Liquidity Maker** has is to *force-close* their channel. However, the operator of an auction venue can request additional information along-side orders before batch execution (as detailed below) to identify premature force closes on-chain in order to penalize the renegading party.

### Channel Lease Funding Protocol

Given the existence of a new modified channel  $C_L$  that supports the channel lease protocol, we require a new way to fund the channels as a channel lease itself will *partially-bind* the following parameters of a given channel:

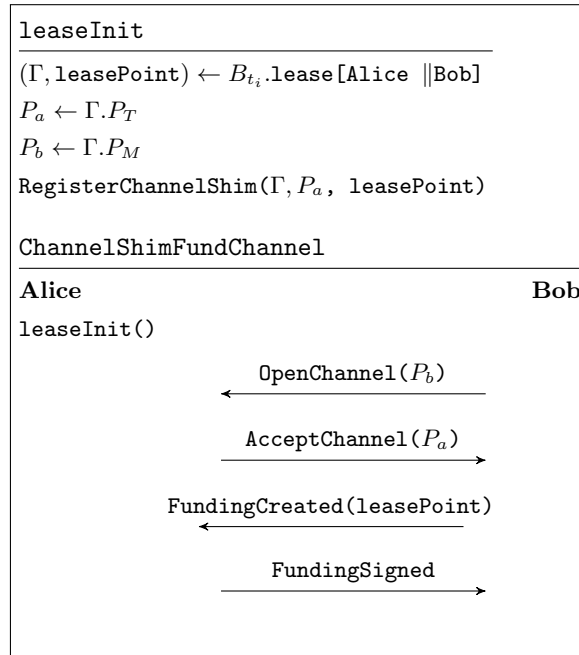
- The size of the channel itself.
- The two multi-sig scripts used within the channel.
- The duration of the channel lease itself.

Rather than modifying the base funding protocol of the Lightning Network, we've opted to instead provide a *new* "side-loadable" partial funding binding API we call a channel shim [26]. A channel shim allows a party  $A$  to *expect* certain parameters of the channel itself. Given the `pendingChanID` which is used to identify an unfinalized channel within the Lightning Network [19], an algorithm

`RegisterChannelShim( $\Gamma, P_a$ , outPoint)` instructs a Lightning Node to use a *prearranged* public key as specified within channel lease  $\Gamma$  for its portion of the multi-sig output used within the ultimate channel.

The existence of the channel shim API allows a node to proceed in the set up of a channel for which it may not yet know the *full* funding transaction to. Given the lease, and the expected outpoint, both sides can fully sign the commitment transaction in a manner that doesn't incur any risk as due to the account structure, they'll need to sign off on the batch funding transaction itself before it can be broadcast.

The funding flow using this new order shim, assuming that match making and market clearing has already occurred, resembles the following flow, with Alice being the **Liquidity Taker** and Bob being the **Liquidity Maker**:



In this modified flow, notice that the taker returns the *pre-generated* public key to use within the multi-sig within the `AcceptChannel` method, and the taker uses the canned `leasePoint` as the input to the funding transaction. As we'll see below the `leasePoint` itself is generated after market clearing, and during the construction of the batch execution transaction. The `leasePoint` as it exists on the batch execution transaction creates a new channel point that satisfies the details of the negotiated channel lease  $\Gamma$ : the amount of the lease, and the pkScript of the two multi-sig keys  $P_a$  and  $P_b$ .

We note that the channel shim abstraction has a number of independent uses including strong layer isolation for more elaborate multi-party channel protocols. Notice how the base **BOLT** funding flow is unmodified, meaning that a higher level application can handle the specified details of the multi-party transaction protocol, while the unmodified Lightning node software manages the underlying channel itself.

#### 6.1.4 Order Structure

With our concrete account structure and channel lease semantics in place, we'll now outline the precise structure of orders as implemented in **Lightning Pool**. As a CLM is a *sealed-bid* auction, the set of active orders within an auction epoch isn't known to a given participant within the auction. Instead, bids are submitted directly to the auctioneer, and may optionally be cancelled between batch epochs as well. Using the fundamental **unit** notion for expressing the *quantity* of an order, we permit partial matching in addition to specifying a *minumum* matchable amount by adding a new field  $M_{match}$  to an order within the  $\Delta_{aux}$  set of additional order attributes.

#### Order Tag Generation & Validation

Next, we specifying order tag generation and validation. Before verifying an order, the order  $\Theta$  itself is serialized in order to generate the message digest of the order. This is done by concatenating each item of the order into a single byte stream, with the set of  $\Delta_{aux}$  attributes specifying a custom serialization mechanism.

For our tag generation, we opt to utilize an SUF-CMA signature scheme. As we target the base Bitcoin blockchain, **Lightning Pool** utilizes Schnorr signatures implemented over the **secp256k1** elliptic curve. In addition to having the trader that submits an order sign the order message digest, we also require that all the backing Lightning nodes of the order also include a signature as well. As the number of backing nodes for a given account may be numerous, rather than accepting individual signatures for each node, we instead require the account operator and all backing nodes to present a *single* Schnorr signature. As we also want have our order tagging scheme be secure against key rogue-key attacks, we select the MuSig2 multi-signature scheme [10].

Given the MuSig2 multi-signature scheme, we define the following algorithms used in our order tagging scheme:

- Let  $\text{Sign}(\{P_0, \dots, P_i\}, M) \rightarrow \sigma$ , be an algorithm that returns a valid MuSig2 multi-signature signed by the set of public keys on message  $M$ .
- Let  $\text{Verify}(\{P_0, \dots, P_i\}, \sigma, M) \rightarrow b$  be an algorithm that returns  $b = 1$  if the passed signature is a valid MuSig2 multi-signature signed by the set of public keys for the message  $M$ .

Given these algorithms, we now define our order tag generation and validation implementations:

$\text{GenOrderTag}(P_{acct}, \Theta)$ <hr/> $m \leftarrow K_{nonce} \  V_{ver} \  P_{acct} \  \Delta_{base} \  \Delta_{aux}$ $tag \leftarrow \text{Sign}(\{\Theta.L_{pub}, \Theta.M_{pub} \dots\}, m)$ $\text{return } tag$
$\text{VerifyOrderTag}(P_{acct}, \Theta)$ <hr/> $m \leftarrow K_{nonce} \  V_{ver} \  P_{acct} \  \Delta_{base} \  \Delta_{aux}$ $b \leftarrow \text{Verify}(\{P_{acct}, P_{M_{pub_0}}, \dots, P_{M_{pub_i}}\}, m)$ $\text{return } b$

We omit the implementations of `SubmitOrder` and `CancelOrder` as they depend on the specific environment in which the auctioneer is implemented in. We only add that the pre-image to an order nonce  $K'_{nonce}$  is known only to the agent that places the order. As we touch on within the future direction section, this commitment structure also has a number of independent uses outside of order cancellation.

### 6.1.5 Node Rating Agencies

Recalling the set of initial requirements that were set out in section 4, it's critical that Lightning Pool reduces information symmetry for the buyer by allowing them to gauge the quality of the node selling a channel lease before they enter into an agreement. In order to achieve this, we introduce the concept of a Node Rating Agency. A rating's agency will allow the buyer of a lease to either query the agency in an ad-hoc manner, or specify that they only wish to be matched with nodes that reside on a certain *tier*.

Let  $T_{node} = \{\hat{t}_0, \dots, \hat{t}_n\}$  be the set of possible ratings that can be given to a Lightning Node with  $\hat{t}_0$  containing *all* known Lightning Nodes, and  $|\hat{t}_i| > |\hat{t}_{i+1}| \forall i < n$ , where  $n$  is the number of available tiers. In other words, we create a series of node sets, with the higher node tiers having less nodes than lower tiers. This creates a natural system of concentric circles where as the tier index increases, the set of nodes shrinks, and eventually only higher quality nodes remain.

Building off this notion of tiered node sets, we define the following algorithm for our node rating agency:

NodeTier( $L_{pub}$ ) <hr/> $\text{nodeMap} \leftarrow \text{map}(T_{node})$ $t_{node} = \text{nodeMap}[L_{pub}]$ <b>return</b> $t_{node}$
---

Given this algorithm, we now specify our of the additional auxiliary order attribute  $\Delta_{aux}$  as:

$$(N_{tier}, \dots) = \Delta_{aux}$$

During match making, we then require the constraint that a given bid order  $\Theta_{bid}$  will only be matched with an ask order  $\Theta_{ask}$  if the following constraint is met:

$$\Theta_{bid}.N_{tier} \geq \text{NodeTier}(\Theta_{ask}.L_{pub})$$

### 6.1.6 Uniform Price Market Clearing & Matching

With our order structure, and the notion of the Node Rating Agency outlined, we now move on to the concrete market clearing and match making within Lightning Pool.

## Order Matching



As mentioned, due to the set of additional constraints outside of simply the posted price and available supply, we employ a multi-attribute match making algorithm. Specifically we opt to utilize a greedy algorithm for the purpose of match making, rather than attempt to find an optimal solution using techniques such as mixed integer linear programming. In this section, we focus primarily on the set of base attributes, leaving the consideration of a trader’s axillary attribute preferences to later work.

First, we define an abstract algorithm which will be used to determine if a given ask order  $\Theta_{ask}$  is compatible preference-wise to a given bid order  $\Theta_{bid}$ :

- **MatchPossible**( $\Theta_{bid}, \Theta_{ask}$ )  $\rightarrow (b, n_{unit})$ . This algorithm returns  $b = 1$  if the given ask and bid are compatible from a match making perspective. If the orders are compatible, then  $n_{unit}$  represents the number of units that can be matched across the two orders.

Given this function we define our implementation of the **MatchMake** algorithm:

<b>MatchMake</b> ( $O: \{\Theta_0, \dots, \Theta_n\}$ ) 1: <b>matchSet</b> $\leftarrow \{\}$ 2: <b>asks</b> $\leftarrow \text{sort}(\text{filter}(O, \Theta_i.O_{type} == Ask))$ 3: <b>bids</b> $\leftarrow \text{sort}(\text{filter}(O, \Theta_i.O_{type} == Bid))$ 4: <b>for</b> $bid$ <b>in</b> $bids$ : 5: <b>for</b> $ask$ <b>in</b> $asks$ : 6: <b>if</b> <b>MatchPossible</b> ( $bid, ask$ ) : 7: $\text{matchSet} = \text{matchSet} \cup (ask, bid)$ 8: <b>return</b> <b>matchSet</b>
---

We note that several optimizations here are possible to reduce the worst-case running time of the algorithm which we leave open for future work. We also assume that the set of valid orders has been filtered out before being passed into this algorithm based on the current target batch fee rate and the posted max batch fee rate of each order.

### Uniform Price Clearing

Once we’ve had our set of candidate matches, we’ll now move on to the market clearing phase. During the market clearing phase, two distinct operations are carried out:

- A trader’s account state is updated to reflect any lease premiums earned due to matches, chain fees paid in the batch execution transaction, fees paid to the auctioneer, and finally the debit for any sold channels from their account.
- We determine the uniform clearing price for a given potential batch.

These two actions comprise the **ClearMarket** and **MarketClearingPrice** algorithms. For our market clearing price, we select the Last Accepted Bid market clearing rule, choosing to go with a buyer’s bid marking price. Given this price clearing algorithm, we now define our implementation of the market clearing algorithms:

$\text{MarketClearingPrice}(\Phi_b)$ <hr style="border: 0; border-top: 1px solid black; margin: 5px 0;"/> <pre> 1: lastPair <math>\leftarrow \Phi_b[\text{len}(\Phi_b) - 1]</math> 2: return lastPair.bid.alpha </pre>
$\text{ClearMarket}(\Psi_A, \Phi_b, \{\Psi_0, \dots, \Psi_n\}, c_{\text{price}})$ <hr style="border: 0; border-top: 1px solid black; margin: 5px 0;"/> <pre> 1: leases <math>\leftarrow \{\}</math> 2: accts <math>\leftarrow \{\}</math> 3: for orderPair in <math>\Phi_b</math> : 4:   lease <math>\leftarrow \text{newLease}(\text{orderPair})</math> 5:   orderPair.taker.balance <math>\mathrel{-=}</math> lease.premium 6:   orderPair.maker.balance <math>\mathrel{+=}</math> lease.premium 7: 8:   orderPair.taker.balance <math>\mathrel{-=}</math> exeFee(lease.amt) 9:   orderPair.maker.balance <math>\mathrel{-=}</math> exeFee(lease.amt) 10:  <math>\Psi_A</math>.balance <math>\mathrel{+=}</math> <math>2 \cdot \text{exeFee}(\text{lease.amt})</math> 11: 12:  leases <math>\leftarrow \text{leases} \cup \text{lease}</math> 13:  accts[orderPair.taker] <math>\leftarrow</math> orderPair.taker 14:  accts[orderPair.maker] <math>\leftarrow</math> orderPair.maker 15: 16: return (<math>\Psi_A</math>, leases, accts) </pre>

Notice that we omit the observance of chain fees, as that will be applied to each input/output during the later batch construction phase.

### 6.1.7 Auction Batch Execution

Once we've cleared the market, we'll now move onto the batch execution phase. During this phase, we'll construct the batch transaction which executes a given batch, and also gather all the necessary witnesses for each participant of the batch so we can properly spend their on-chain account outputs. Remember that due to the non-custodial structure, a trader will fully validate a given batch before they sign off on it.

#### Batch Transaction Construction

A given batch transaction contains the following inputs: the set of trader account inputs involved in the batch, and the input of the master auctioneer itself. In addition to these inputs which can only be spent each each user authorized the proposed batch, we also add the following outputs: a trader's new incremented account output which reflects the market clearing, the set of channels created as part of the channel lease, and the incremented auctioneer account. Note that the format of the batch transaction itself may change multiple times during this phase if participants reject the batch, or if fee changes causing the auctioneer to consider a subset of the prior set of orders.

Taking this into consideration, we define our implementation of **Construct-Batch** as follows:

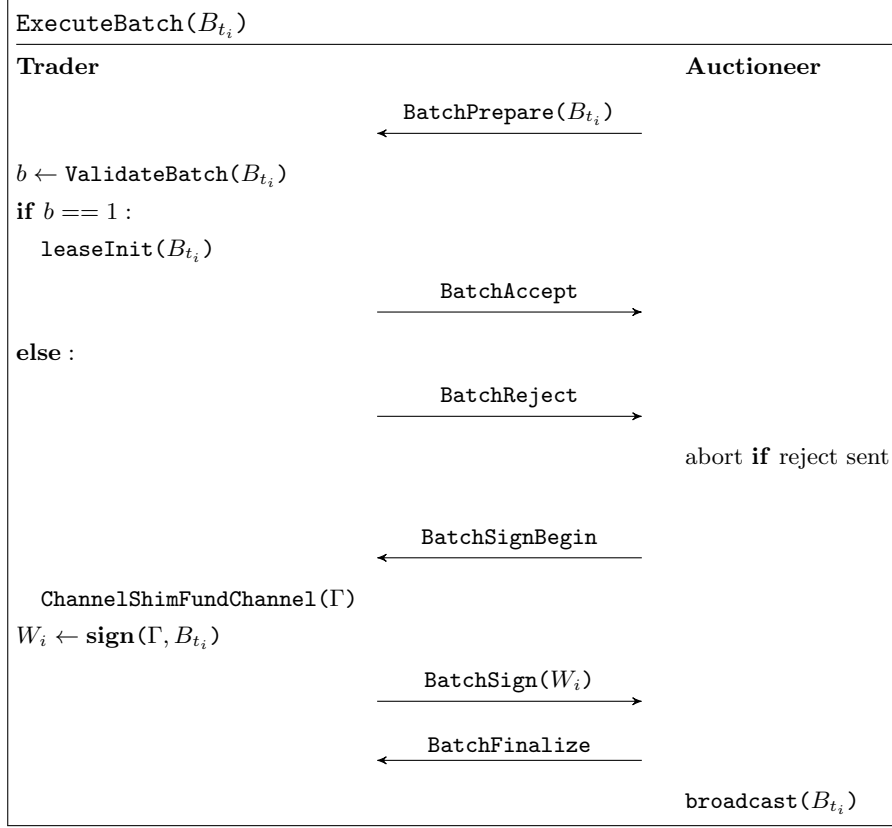
ConstructBatch( $\Delta_i$ )	
1 :	<b>tx</b> $\leftarrow$ <b>newTx</b> ()
2 :	<b>tx.addIn</b> ( $\Delta_i.\Psi_A$ )
3 :	
4 :	<b>for</b> <i>acct</i> <b>in</b> $\Delta_i.\Psi$ :
5 :	<b>tx.addIn</b> ( <i>acct.prevOut</i> )
6 :	<b>for</b> <i>acct</i> <b>in</b> $\Delta_i.\Psi'$ :
7 :	<i>acct.Value</i> $\mathrel{--}$ <b>feeShare</b> ( $\Delta_i$ , <i>acct</i> )
8 :	<b>tx.addOut</b> ( <i>acct.txOut</i> )
9 :	
10 :	<b>for</b> <i>lease</i> <b>in</b> $\Delta_i.\Gamma$ :
11 :	<b>tx.addOut</b> ( <i>lease.txOut</i> )
12 :	
13 :	<i><math>\Psi_A.value</math></i> $\mathrel{--}$ <b>feeShare</b> ( $\Delta_i$ , $\Psi_A$ )
14 :	<b>tx.addOut</b> ( $\Delta_i.\Psi'_A$ )
15 :	<b>return</b> <i>tx</i>

### Batch Transaction Execution

Once the batch has been constructed, the auctioneer then needs to propose the batch to each trader, and collect the necessary set of signatures required to spend each trader's account input. During this execution phase, we assume that this is the final set of traders that wish to be a part of this batch. Once the auctioneer has all the necessary signatures to broadcast a batch, the batch execution transaction can be broadcast, ending this auction epoch.

Batch transaction execution itself is a multi-party protocol wherein the auctioneer presents a valid auction batch to all involved parties, the parties verify the batch, before finally signing off by resending valid signatures of the batch (thereby attesting to it) which are require to execute the batch by committing it in the Bitcoin blockchain.

We define batch execution in the context of Lightning Pool as follows:



Once the auctioneer has gathered all the signatures  $W_i$  from each participant of the batch, then it can sign its own input in the batch transaction, and broadcast the batch thereby finalizing execution. We now provide additional insight with respect to the meaning of each of the messages sent above.

- The **BatchPrepare** message kicks off batch execution and is used to *propose* a new potential auction batch to an end user client. Upon receipt of this message, the client will attempt to verify the batch using the **ValidateBatch** algorithm.
- If the batch is *invalid* from the client's PoV, then a **BatchReject** message is sent. Upon receipt of the reject message the auctioneer will resume the protocol, excluding the rejecting client.
- Otherwise, the client accepts the batch after verifying its integrity by sending the **BatchAccept** message.
- In order to synchronize the creation of funding shims and the channel funding protocol itself, the auctioneer send the **BatchSignBegin** message. After sending the reject, the *taker* will register a funding shim to ensure it's able to properly handle the incoming funding request in the next phase.
- Upon receipt of the **BatchSignBegin** message, the client then is ready to fund the channel, using the **leasePoint** which is present on the batch execution transaction  $B_{t_i}$ .

- Once the funding is complete (both sides have a valid commitment transaction), then the trader will sign its input to present a valid signature to the auctioneer, so the batch transaction can be completed.
- The final set for the auctioneer is to commit the new batch to disk, and broadcast the transaction to the Bitcoin network.

#### 6.1.8 Sidecar Channel Market Clearing & Batch Execution

As mentioned in the background section, a Channel Lease Marketplace is able to also implement the abstract notion of a "sidecar channel". We begin by first defining a sidecar channel as so:

**Definition 6.1.** (Sidecar Channel. A **Sidecar Channel** is defined as  $\Sigma_c = \{A_r, B_s, C_g, N_{sat_{in}}, N_{sat_{out}}\}$ , where:

- $A_r$  is the *receiver* of a new channel which may contain both inbound and outbound liquidity. This party may or may not already be a participant in the CLM.
- $B_s$  is the seller of a normal channel lease within the CLM, and holds an active trading account.
- $C_g$  is the *gifter* of a sidecar channel who wishes to onboard  $A_r$  to the Lightning Network. Note that  $A_r$  may not even have any Bitcoin at all.
- $N_{sat_{in}}$  is the number of Bitcoin expressed in satoshis of the channel itself, which will be available as inbound channel bandwidth to  $A_r$ .
- $N_{sat_{out}}$  is the number of Bitcoin which is to be pushed [19] (`push.amt`) to the receiver of the side car channel, allowing them to both send and receive.

At a high level, a sidecar channel allows Alice to buy a channel for Carol via Bob. Pool implements sidecar channels by making a series of small modifications to the normal marketplace operations. An end-to-end workflow resembles the following:

- Alice who already has an account in the CLM places an order to buy a channel. However, rather than specifying the public key and connection details of her own node, she uses Carol's information instead.
- During market clearing, if  $N_{sat_{out}} > 0$ , then Alice will pay an *additional* amount of  $N_{sat_{out}}$  satoshis directly into the account of Bob.
- During batch execution, rather than Alice registering a channel shim, she informs Carol to do so in an expectation step. Bob then uses the new  $N_{sat_{out}}$  satoshis in his account to push the Bitcoin over to the channel he created between himself and Alice.
- To ensure the agreement is upheld, Carol uses a channel acceptor [27] predicate to validate that the proper amount has been pushed.

Aside from the above modifications to our market clearing and batch execution, everything else remains untouched. The end result is the ability to purchase a channel for a 3rd party in the network in a trust-minimized manner. This new ability is akin to being able to insert an edge in the network (a new channel) between any two arbitrary but cooperating parties. As an example, an exchange could use this flow to allow a client to withdraw directly into a *new* channel.

### 6.1.9 LSAT as Pool Tickets

As an added layer against spam and resource exhaustion attacks, Lightning Pool uses Lightning Service Authentication Tokens, or LSATs [7] when interacting with all users. In order to perform operations such as creating an account, querying batch snapshots, etc, a valid LSAT is required. The auctioneer is able to dynamically raise the price of an LSAT which is expressed in satoshis if anomalous behavior is detected.

In addition to using LSATs to throttle or rate limit clients, the auctioneer is also able to use them as a mechanism to offer up historical chain data for sale to 3rd party observers (those without active trading accounts) of the system.

## 6.2 The Lightning Pool Shadowchain

In this section, we complete the Lightning Pool system by demonstrating out its implementation of a Channel Lease Marketplace can be implemented using our shadowchain application overlay framework.

### 6.2.1 Lightning Pool Accounts as Lifted UTXOs

First, we link the concept of our non-custodial accounts in the CLM realm to a lifted UTXO. The process of lifting and unlifting a UTXO is simply a series of operations required to create, modify or close an account:

<b>LiftUTXO</b> ( $T_{expiry}, \{U_{N_0}, \dots, U_{N_n}\}, P_0$ )
1: <b>inputs</b> $\leftarrow \{U_{N_0}, \dots, U_{N_n}\}$ 2: <b>return</b> <b>NewAccount</b> ( $T_{expiry}, P_0, \text{inputs}$ )
<b>UnliftUTXO</b> ( $\Psi_U$ )
1: $(b, ) \leftarrow \text{ModifyAccount}(\Psi_U, P_{auction_p})$ 2: <b>return</b> $b$
<b>ExitChain</b> ( $\phi_U, B_{height}$ )
1: <b>if</b> $\phi_U.T_{blocks} < B_{height}$ : 2: <b>return</b> 0 3: $(b, ) \leftarrow \text{ModifyAccount}(\Psi_U, P_{auction_p})$ 4: <b>return</b> $b$

Note that it's also possible to upgrade Lifted UTXOs as implemented within Pool, as a given user is able to use the latest features in Bitcoin script to achieve the same functionality. However, due to privacy implications, it may

be preferred to have all account scripts, and further all scripts within a batch execution transaction be uniform.

### 6.2.2 Auction Batch Proposal

Next, we move unto patch proposal and acceptance. An auction batch within a CLM maps 1:1 to the concept of blocks in the shadowchain domain. Given this insight, we now define the **ConstructBlock** and **ProposeBlock** algorithms:

<b>ConstructBlock</b> ( $\phi_{live}, T_{xn}, E_{exe}, \Delta_F$ )	<b>ProposeBlock</b> ( $B_S, \phi_{live}$ )
$\Phi_b \leftarrow \text{MatchMake}(T_{xn})$	$b \leftarrow \text{ValidateBatch}(B_S, \phi_{live})$
$c_{price} \leftarrow \text{MarketClearingPrice}(\Phi_b)$	<b>return</b> $b$
$\text{ClearMarket} \leftarrow \Delta_F$	
$\Delta_i \leftarrow \text{ClearMarket}(\Psi_A, \Phi_b, \Psi, c_{price})$	
<b>return</b> $\Delta_i$	

The process of constructing a new block walks through each of the phases within the auction itself. Note that a block in our system can only be constructed if there exist a valid market clearing given the set of orders ( application-specific transactions). The auctioneer then proposes the block to each party within the **BatchPrepare** message. Traders are then free to accept or reject a given block.

### 6.2.3 Shadowchain Batch Execution

Now that we're able to lift/unlift UTXOs and propose blocks within our Shadowchain, we now define the series of methods that will be utilized to allow clients to execute the their local version of the state transition function to accept new proposed batches:

<b>CommitBlock</b> ( $B_S$ )
1 : $(b, TX_{id}) \leftarrow \text{ExecuteBatch}(B_S)$
2 : <b>return</b> $b$

The **ExecuteBatch** algorithm does most of the heavy lifting here. Note that if for whatever reason, execution fails, then  $b = 0$  is returned, and we resume our normal state machine execution loop. At the end of the **CommitBlock** algorithm all participants have the latest block (as they need to be given the block in order to sign off on it), and the block is broadcast to the Bitcoin blockchain. After this phase, it's possible to continue committing new blocks without waiting for prior blocks to confirm. Due to this flexibility, the Orchestrator of the Lightning Pool shadow chain is then able to *optimistically* perform transaction cut-through to combine several logical blocks into a single Bitcoin transaction.

### 6.2.4 Unconfirmed Batch Cut-Through

Recall that a shadowchain block can also be optimistically aggregated into a single block. In the domain of the Lightning Pool shadowchain, combining blocks requires ensuring that all produced channel leases will still exist in the final combined blocks, and the end state of each account remelts any consecutive market clearing opportunities:

<p><b>CoalesceBlocks</b>(<math>F: \{B_{S_0}, \dots, B_{S_N}\}</math>)</p> <pre> 1: <b>inputs</b> <math>\leftarrow \text{set}(F)</math> 2: <b>leases</b> <math>\leftarrow \text{extractLeases}(F)</math> 3: <b>accts</b> <math>\leftarrow \text{endAcctState}(F)</math> 4: 5: <b>tx</b> <math>\leftarrow \text{ConstructBatch}(\Delta_i(\text{inputs}, \text{leases}, \text{accts}))</math> 6: <b>b</b> <math>\leftarrow \text{ExecuteBatch}(t)</math> 7: <b>return</b> <i>b</i> </pre>
--

For brevity, we omit the referenced internal algorithms, however the naming is intended to be intuitive. Given a series of blocks  $F$ , we need to extract all the inputs referenced in each block (the series of accounts), extract the set of all leases created within the block, and also the ending account state of all accounts involved in the prior blocks. Note that one account may participate in all or some of the prior blocks. The process of block cut-through allows us to only manifest the *ending state* of their account and elide the intermediate states from the PoV of the blockchain. With these summaries constructed, the auctioneer then construct a new batch, and executes it using the normal algorithms.

Scalability gains are had by only manifesting the final state of each account, as well as removing the need to manifest all intermediate transaction within the blockchain. All participants have an incentive to participate in this optimistic block cut-through as they'll also end up paying less chain fees as they only need to pay for their account input+output, and any leases created *once*. If any of the intermediate blocks confirm instead of the cut-through block, then the process can be repeated with the new set of unconfirmed blocks.

### 6.2.5 Auction Upgrades

Finally, we define the process by which we upgrade the CLM shadowchain itself. As this is an off-chain process, each participant of the shadowchain is able to execute the **UpgradeChain** algorithm simply by updating their end client software. In addition to this we utilize two upgrade extension points:

- The version of a given order.
- The **batchVersion** sent within the **BatchPrepare** message.

The version in each order allows us to add new order types over time which implement new match making related preference expression, or brand new channel types. The version sent along during batch execution allows us to modify attributes such as the fee sharing scheme, or the structure of the batch execution transaction itself.

## 7 Security Analysis

Similar to Lightning Loop, the Lightning Pool backend server is may be closed source, but clients are able to fully verify and audit each phase of the auction. At a high level, Pool can be seen as a “shadow chain” anchored in the base Bitcoin blockchain. The shadow chain validates modifications to a subset of the



UTXO set (the Pool accounts) with the auctioneer acting as a block proposer to extend the chain. State transitions are validated and accepted by those that are involved in a new block (the auction batch). Newer clients are even able to audit the prior history of the system in order to ensure proper operation. Pool uses the Bitcoin blockchain for what it’s best for: global censorship resistant batch execution.

Leveraging this shadow chain structure, users remain in control of their funds at all times, and will only enter into agreements that they’re able to fully verify, ensuring that channels are properly constructed and that the market is operating as expected. Compared to existing centralized exchanges with off-chain order execution, Pool has a number of attractive security properties:

- As a non-custodial system, users are in control of their funds at all times.
- A purchased LCL will result in the creation of a channel with parameters that capture the preferences expressed in the initial order.
- If the auctioneer server is hacked, the breach doesn’t unilaterally compromise user funds.
- Orders by one trader cannot be used to spoof orders by another trader.
- Clients are able to verify and audit all operations carried out by the server during batch construction including proper order matching.

## 8 Future Directions

As is clients verify each Shadowchain blocks within the CLM system, but they have no assurance that their order was actually included in the mach making function. In this manner, the auctioneer can silently ignore a set of orders. To remedy this, it may be possible for the auctioneer to publish an order transparency authenticated data structure to give users a merkle leaf receipt of proper order inclusion. Agents in the marketplace would then exchange their subjective order tree roots to confirm all orders have been included before they enter the batch execution phase.

Rather than sending over a full block, the auctioneer can instead send over a zero-knowledge argument of proper block validity. This would improve the privacy of the system, as less information about the underlying order book is leaked to participants. The proof would be rooted at the order merkle tree root in order to give traders more assurance that their orders were included.

As is, the maker re receives their premium all at once. However it may be possible to set up another uni-directional channel in order to stream the interest in real-time. We call this concept of a uni-directional channel used to stream the owed interest with each new block a coupon channel. It may further be possible to allow others to buy/sell the future cash flows of the ”coupon channels”.

## 9 Related Work

The Celar Network [33] puts forth a concept of A *Liquidity Backing Auction*. Their auction differs from ours in that they opt to utilize a Vickery-Clarke-Groves auction, and unlike our double-call auction, the auction described in the

system takes place with a designated on-chain Ethereum contract with a single seller. Rather than paying out interest in the currency of the base blockchain, the Celar network opts to instead issue a new sub-currency which is used to reward liquidity providers, which actually an IOU, thus being subject to default risk. In addition, all order submission and auction clearing takes place on-chain rather than off-chain as implemented with Lightning Pool.

The Gnosis Prediction [34] market uses an on-chain batched auction to provide liquidity for multiple distinct assets. Compared to our work, all operations happen on-chain rather than off-chain, and they opt to use mixed-integer linear programming to derive an optional solution for their market clearing target.

Independent Bitcoin developer ZmnSCPxj’s concept of *Smart Contracts Unchained* [31] is similar to our Shadowchain overlay application framework. However, their system is based on a concept of escrows which involved a 3rd party, whereas our concept of Lifted UTXOs only concerns the user and the Shadowchain orchestrator. In addition to this, they implement their system with a single output that’s used by all participants. Shadowchains on the other hand give each participant an individual UTXO, which allows for a more flexible participant set in proposed batches. Finally, their system may allow for funds to be stolen with the collusion of one of the participants, and the smart contract platform service. Lifted UTXOs within the context of shadowchains are fully non-custodial.

Lisa Neigut, and Casey Rodamor put forth an idea to extend the Lightning p2p network to support liquidity advertisements for peers seeking to solicit new dual funded channels [28]. Compared to our work, their proposal is an adhoc p2p bulletin board without a true venue for market discovery. We also target single funded channels as we note that it’s possible to simulate dual funded channels within our auction using a pair of bid and ask orders matched between the same parties.

Bitrefill’s Thor [25] service is similar to our work, however there exist only a single seller (Bitrefill) which names a price (no preference incorporation) and doesn’t guarantee a lifetime of the channel.

## 10 Conclusion

In this work, we’ve put forth a new abstraction over capital obligations in the Lightning Network, which we call channel leases. Channel leases can be bought and sold on a Channel Liquidity Marketplace directly solving a series of bootstrapping challenges the network faces, commonly referred to as: the inbound liquidity problem. To implement a Channel Lease Marketplace in a secure manner, we put forth the concept of the Shadowchain application framework which is of independent use. We then concretely construct Lightning Pool, the first CLM implemented on top of the base Bitcoin blockchain. Lightning Pool allows those with idle capital to earn yield on their Bitcoin, and also allows those that need inbound to receive over the network to obtain a reliable source of incoming payment bandwidth. The exigence of systems such as Lightning Pool creates new revenue sources for routing nodes on the network, and also creates a new form of yield bearing instrument for the greater Bitcoin ecosystem.

## References

- [1] Joseph Poon & Tadge Dryja, *The Bitcoin Lightning Network: Scalable Off-Chain Instant Payments*, DRAFT Version 0.5.9.2, <https://lightning.network/lightning-network-paper.pdf>, 2016
- [2] Lewis Gudgeon, Pedro Moreno-Sanchez, Stefanie Roos, Patrick McCorry, Arthur Gervais, *SoK: Layer-Two Blockchain Protocols*, <https://eprint.iacr.org/2019/360.pdf>, 2019
- [3] Department of the Treasury, *Uniform-Price Auctions Update of the Treasury Experience*, <https://www.treasury.gov/press-center/press-releases/Documents/upas.pdf>, Oct 1998
- [4] Lightning Network Developers, *Anchor outputs*, <https://github.com/lightningnetwork/lightning-rfc/pull/688>, Oct 28, 2019
- [5] Simina Brânzei, Erel Segal-Halevi, Aviv Zohar, *How to Charge Lightning*, <https://arxiv.org/abs/1712.10222>, Dec 27, 2017
- [6] Paul Milgrom & Robert J. Weber, *The Value of Information in a Sealed-Bid Auction*, <http://web.stanford.edu/~milgrom/publishedarticles/ValueOfInformation.pdf>, 1982
- [7] Olaoluwa Osuntokun, *LSAT: Authentication and Payments for the Lightning-Native Web*, <https://lightning.engineering/posts/2020-03-30-lsat/>, March 30, 2020
- [8] Olaoluwa Osuntokun & Conner Fromknecht, *AMP: Atomic Multi-Path Payments over Lightning*, <https://lists.linuxfoundation.org/pipermail/lightning-dev/2018-February/000993.html>, Feb 6, 2018
- [9] Thomas Ristenpart and Scott Yilek, *The Power of Proofs-of-Possession: Securing Multiparty Signatures against Rogue-Key Attacks* <https://www.iacr.org/archive/eurocrypt2007/45150228/45150228.pdf>,
- [10] Jonas Nick, Tim Ruffing, and Yannick Seur *MuSig2: Simple Two-Round Schnorr Multi-Signatures*, <https://eprint.iacr.org/2020/1261.pdf> October 11, 2020
- [11] SIMON PARSONS, JUAN A. RODRIGUEZ-AGUILAR, MARK KLEIN, *Auctions and bidding: A guide for computer scientists*, <http://www.sci.brooklyn.cuny.edu/~parsons/courses/7412-fall-2011/papers/bluffers-final.pdf>, 2020
- [12] Meni Rosenfeld, *Analysis of Bitcoin Pooled Mining Reward Systems*, <https://arxiv.org/abs/1112.4980>, Dec 21, 2011
- [13] Fidelity Bonds, [https://en.bitcoinwiki.org/wiki/Fidelity\\_bonds](https://en.bitcoinwiki.org/wiki/Fidelity_bonds), 2020
- [14] Greg Maxwell, *Transaction cut-through*, <https://bitcointalk.org/index.php?topic=281848.0>, Aug 26, 2013

- [15] Peter Cramton, *Innovation and Market Design*, <http://www.cramton.umd.edu/papers2005-2009/cramton-innovation-and-market-design.pdf>, May 9, 2008
- [16] Peter Cramton, *Market Design: Harnessing Market Methods to Improve Resource Allocation*, <http://www.cramton.umd.edu/papers2010-2014/cramton-market-design.pdf>, Oct 16, 2010
- [17] Thomas Kerin & Mark Friedenbach, *BIP 113*, <https://github.com/bitcoin/bips/blob/master/bip-0113.mediawiki>, Aug 10, 2016
- [18] gmaxwell, *Transaction cut-through*, <https://bitcointalk.org/index.php?topic=281848.0>, Aug 10, 2016
- [19] Lightning Network Developers, *BOLT #2: Peer Protocol for Channel Management*, <https://github.com/lightningnetwork/lightning-rfc/blob/master/02-peer-protocol.md>,
- [20] Lightning Network Developers, *BOLT #4: Onion Routing Protocol*, <https://github.com/lightningnetwork/lightning-rfc/blob/master/04-onion-routing.md>,
- [21] Lightning Network Developers, *BOLT #3: Bitcoin Transaction and Script Formats*, <https://github.com/lightningnetwork/lightning-rfc/blob/master/03-transactions.md>,
- [22] Lightning Network Developers, *BOLT #7: P2P Node and Channel Discovery*, <https://github.com/lightningnetwork/lightning-rfc/blob/master/07-routing-gossip.md>,
- [23] Alfred E. Kahn, Peter C. Cramton, Robert H. Porter, Richard D. Tabors *Uniform Pricing or Pay-as-Bid Pricing: A Dilemma for California and Beyond*, <ftp://cramton.umd.edu/papers2000-2004/kahn-cramton-porter-tabors-uniform-or-pay-as-bid-pricing-ej.pdf>, 2001
- [24] Eric Budish, Peter Cramton, John Shim, *THE HIGH-FREQUENCY TRADING ARMS RACE: FREQUENT BATCH AUCTIONS AS A MARKET DESIGN RESPONSE*, <https://faculty.chicagobooth.edu/eric.budish/research/HFT-FrequentBatchAuctions.pdf>, November 2015
- [25] Bitrefill, *Thor: Lightning Channel-Opening Service*, <https://www.bitrefill.com/thor-lightning-network-channels/?hl=en>, Jan 9, 2019
- [26] Olaoluwa Osunoktun, *chanfunding: create new package to abstract over funding workflows*, <https://github.com/lightningnetwork/lnd/pull/3659>, Oct 31, 2019
- [27] Eugene Siegel, *rpc: bi-directional streaming for predicate-based channel acceptance*, <https://github.com/lightningnetwork/lnd/pull/3039>, May 3, 2019
- [28] Lisa Neigut & Casey Rodamor, *Proposal for Advertising Channel Liquidity*, <https://lists.linuxfoundation.org/pipermail/lightning-dev/2018-November/001532.html>, Nov 7, 2018

- [29] Peter Todd, *BIP 65*, <https://github.com/bitcoin/bips/blob/master/bip-0065.mediawiki>, Oct 10, 2014
- [30] ZmnSCPxj, *Towards a Market for Liquidity Providers – Enforcing Minimum Channel Lifetime*, <https://lists.linuxfoundation.org/pipermail/lightning-dev/2018-November/001555.html>, Nov 10, 2018
- [31] ZmnSCPxj, *Smart Contracts Unchained*, <https://zmnsctxj.github.io/bitcoin/unchained.html>, October 2020
- [32] Oğuzhan Ersoy, Stefanie Roos and Zekeriya Erkin, *How to profit from payments channels*, <https://arxiv.org/pdf/1911.08803.pdf>, October 2020
- [33] ScaleSphere Foundation Ltd. ("Foundation"), *Celer Network: Bring Internet Scale to Every Blockchain*, <https://www.celer.network/doc/CelerNetwork-Whitepaper.pdf>, June 26, 2018
- [34] Tom Walther, *Multi-token Batch Auctions with Uniform Clearing Prices*, <https://github.com/gnosis/dex-research/blob/master/BatchAuctionOptimization/batchauctions.pdf>, September 3, 2020